

Group Assignment Trimester 1 25/26

2025-09-16

Task

Description

You are a consultant for Pacific Insights Analytics (PIA), a boutique data analytics firm specializing in workforce performance. PIA has been hired by Huaxia Communications, a major Chinese company operating several large call centers across the country. Huaxia is facing rising operational challenges. Staff attrition is driving up recruitment and training costs and retaining top performers is becoming increasingly difficult. To address these issues, Huaxia has provided you with historical data from one of its call centers. This data comes from multiple administrative sources, including worker surveys, HR records, and call center productivity logs. Because the data was collected from different systems, it contains inconsistencies, formatting problems, and other quality issues. Your task is to clean, explore, and analyze the data in order to generate clear, evidence-based recommendations that can help Huaxia's management improve retention and boost productivity. Important: You should not run any regression models. The goal is to understand the data and identify patterns that can inform management decisions through EDA and hypothesis testing.

Declaration of AI Usage

In the preparation of this project, we made use of AI tools, specifically OpenAI's ChatGPT, as a support resource. The tool was employed to enhance efficiency in tasks such as structuring ideas, clarifying concepts, refining wording, and checking technical accuracy.

At all times, the use of AI served as a complement to my own knowledge, analysis, and critical thinking. The outputs generated were carefully reviewed, adapted, and integrated by us to ensure originality, accuracy, and academic integrity. The final work reflects my own understanding, reasoning, and skills, with AI functioning only as a leverage of these.

Setup

Packages

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.2      v tibble    3.3.0
## v lubridate  1.9.4      v tidyr     1.3.1
## v purrr      1.1.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```

library(haven)
library(dplyr)
library(ggplot2)
library(ggcorrplot)
library(ISOweek)
library(stargazer)

##
## Please cite as:
##
## Hlavac, Marek (2022). stargazer: Well-Formatted Regression and Summary Statistics Tables.
## R package version 5.2.3. https://CRAN.R-project.org/package=stargazer
library(fmsb)

```

Loading the data

```

data_perf <- read_dta("../data/performance_labeled.dta")
data_attitude <- read_dta("../data/attitude_labeled.dta")
data_endperiod <- read_dta("../data/endperiod_outcomes_labeled.dta")
data_sum_vol <- read_dta("../data/summary_volunteer_labeled.dta")
data_wage_new <- read_dta("../data/wage_new_labeled.dta")

# convert to tibbles
performance <- as_tibble(data_perf)
attitude <- as_tibble(data_attitude)
endperiod <- as_tibble(data_endperiod)
volunteer <- as_tibble(data_sum_vol)
wage <- as_tibble(data_wage_new)

# remove raw data
rm(data_perf)
rm(data_attitude)
rm(data_endperiod)
rm(data_sum_vol)
rm(data_wage_new)

```

Cleaning data

Handy functions for cleaning

Here come some handy functions to reduce duplicate code for cleaning the data:

```

# convert to numeric, if its a number (not NA)
as_numeric <- function (x) {
  x <- as.character(x)
  # only convert if it looks numeric
  if (grepl("^[-+]?[0-9]*\\.?[0-9]+([eE][-+]?[0-9]+)?$", x)) {
    as.numeric(x)
  } else {
    NA_real_
  }
}

```

```

# convert to integer, if its a number (not NA)
as_integer <- function(x) {
  # only convert if it looks numeric
  x <- as.character(x)
  if (grepl("^[~+]?[0-9]*\\.?[0-9]+([eE][~+]?[0-9]+)?$", x)) {
    as.integer(x)
  } else {
    NA_real_
  }
}

# Extract a numeric values correctly from a string value, considering multiple cases
extract_numeric <- function(x) {
  # remove leading/trailing spaces
  x <- trimws(x)

  # normalize dash types (en-dash, em-dash) to minus
  x <- gsub("-", "-", x)

  # allow scientific notation: replace commas in mantissa
  if (grepl("[eE]", x)) {
    # if comma before e → decimal comma
    x <- sub(",", ".", x)
    # remove thousand separators before e
    x <- gsub("\\.(?=[0-9]{3}[eE])", "", x, perl = TRUE)
    x <- gsub(",(?=[0-9]{3}[eE])", "", x, perl = TRUE)
    return(as.numeric(x))
  }

  # remove spaces
  x <- gsub(" ", "", x)

  # case 1: comma used as decimal
  if (grepl(",", x) & !grepl("\\.", x)) {
    x <- gsub(",", ".", x)
  }

  # case 2: both comma and dot appear
  # assume comma is thousands separator, remove it
  if (grepl(",.*\\.", x)) {
    x <- gsub(",", "", x)
  }

  as.numeric(x)
}

# find all non numeric characters in a string and return them
find_non_numeric_chars <- function(x) {
  non_numeric <- gsub("[0-9]", "", x) # remove digits, minus, dot, comma
  unique(non_numeric[non_numeric != ""])
}

# find all rows with non numeric characters in a string and return them
find_non_numeric_rows <- function(x) {
  non_numeric <- !grepl("^[~+]?[0-9]+([0-9]+)?$", x) # remove digits, minus, dot, comma

```

```

non_numeric <- non_numeric[non_numeric != ""]
x[non_numeric]
}

```

Clean volunteer

```
glimpse(volunteer)
```

```

## Rows: 140
## Columns: 9
## $ personid <chr> "33350", "40034", "31292", "21654", "36908", "13980", "44794~
## $ age <dbl> 26, 19, 27, 29, 23, 1, 23, 25, 21, 23, 23, 30, 22, 22, 23, 2~
## $ tenure <dbl> 22, 9, 24, 37, 15, 51, 3, 42, 3, 34, 24, 25, 3, 66, 9, 70, 4~
## $ children <chr> "0", "0", "0", "0", "0", "0", "0", "0", "0", "0", "0", "0", ~
## $ bedroom <chr> "1", "1", "1", "1", "1", "1", "1", "1", "1", "1", "1", "1", ~
## $ commute <chr> "120", "300", "135", "80", "60", " 120", "40 min", "60", "1~
## $ gender <chr> "m", "woman", "F", "f", "m", "woman", "male", "woman", "MALE~
## $ married <chr> "SINGLE", "N", "NO", "N", "NO", " single", " N", "NO", "sin~
## $ high_educ <chr> "Y", "no", "YES", "YES", "N", "NO", "n", "YES", "no", "Y", "~

```

```
stargazer(as.data.frame(volunteer), type = "text", median = TRUE, header = FALSE)
```

```

##
## =====
## Statistic N    Mean St. Dev.  Min  Median   Max
## -----
## age      140 23.536  4.352    1    23    35
## tenure   140 22.971 24.296  2.000 13.000 150.000
## -----

```

```

# check character columns with numerical values to extract for non numeric chars to
# be aware for special cases

```

```
find_non_numeric_chars(volunteer$commute)
```

```

## [1] " "      " min"    " MIN"    " "      " minute" " minutes" ". h"
## [8] " MIN"

```

```
find_non_numeric_rows(volunteer$commute)
```

```

## [1] " 120"      "40 min"      " 50 MIN"      " 170"      " 210"
## [6] "40 minute" "240 minutes" "1.2 h"        " 180"      " 90 MIN"
## [11] "40 minute" "30 minutes"  " 180"        "120 MIN"   " 120"
## [16] "50 min"    "100 min"     " 120"        "60 minutes" "1.2 h"
## [21] "180 minutes" " 150"       "2.3 h"

```

```
# here we only need to take care of the decimal dots (will be added to the extraction function)
```

We can see a lot of columns with numerical values within strings and additional chars. Furthermore, categorical variables like married and gender contain different labels for the same categories.

```

# convert to integer to remove leading zeros an than to str
volunteer$personid <- as.character(as.integer(volunteer$personid))
# Clean Numbers with extra characters
# Remove non-numeric characters, except dots and commas and convert to numeric
volunteer$commute <- sapply(volunteer$commute, extract_numeric)

```

```
## Warning in FUN(X[[i]], ...): NAs introduced by coercion
```

```

## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion

volunteer$commute <- sapply(volunteer$commute, as_numeric)

# convert "no" strings in the children column to 0 children
volunteer$children <- ifelse(volunteer$children %in% c("No", ""), "0", volunteer$children)
# convert to boolean by first converting to integer
volunteer$children <- as.logical(sapply(volunteer$children, as_integer))
# Convert Tenure to integer, as float dont make sense for months
volunteer$tenure <- sapply(volunteer$tenure, as_integer)

# convert bedroom to boolean by first converting 0s to False, 1 to true
volunteer$bedroom<- ifelse(
  volunteer$bedroom == "0",
  0,
  1
)
volunteer$bedroom <- sapply(volunteer$bedroom, as_integer)
# normalize gender data by
# lower all charcters
volunteer$gender <- tolower(volunteer$gender)
# trim whitespaces
volunteer$gender <- trimws(volunteer$gender)
# cast all possible options for male and female
volunteer$gender <- ifelse(volunteer$gender %in% c("m", "male", "man"), "male", volunteer$gender)
volunteer$gender <- ifelse(volunteer$gender %in% c("woman", "female", "f", "w","fem"), "female", volunteer$gender)

# using a similar approach to the gender we convert the married status now
volunteer$married <- tolower(volunteer$married)
# trim whitespaces
volunteer$married <- trimws(volunteer$married)
# cast all possible options for married/not married
volunteer$married <- ifelse(volunteer$married %in% c("married", "yes", "y", "true"), "TRUE", volunteer$married)
volunteer$married <- ifelse(volunteer$married %in% c("no", "not married", "n", "na", "single","false"), "FALSE", volunteer$married)

# using a similar approach to the married column we convert the high_educ status now
volunteer$high_educ <- tolower(volunteer$high_educ)
# trim whitespaces
volunteer$high_educ <- trimws(volunteer$high_educ)
# cast all possible options for high_educ/not high_educ
volunteer$high_educ <- ifelse(volunteer$high_educ %in% c("yes", "y"), "TRUE", volunteer$high_educ)
volunteer$high_educ <- ifelse(volunteer$high_educ %in% c("no", "n","false"), "FALSE", volunteer$high_educ)

# convert to boolean
volunteer$high_educ <- as.logical(volunteer$high_educ)

# create numerical variable cols for gender and married
volunteer$gender_num <- sapply(ifelse(volunteer$gender=="male",0,1), as_integer)
volunteer$married_num <- sapply(ifelse(volunteer$married==FALSE,0,1), as_integer)
volunteer$high_educ_num <- sapply(ifelse(volunteer$high_educ==FALSE,0,1), as_integer)

```

```
volunteer$children_num <-sapply(ifelse(volunteer$children==FALSE,0,1), as_integer)
```

```
# remove duplicates on same personid
volunteer <- volunteer %>%
  distinct(personid, .keep_all = TRUE)
```

```
# Now we have much better insights using stargazer
stargazer(as.data.frame(volunteer), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N    Mean    St. Dev. Min Median Max
## -----
## age            135 23.452    4.315    1    23    35
## tenure         135 23.385    24.608    2    13   150
## children       135  0.148    0.357    0     0     1
## bedroom        135  0.978    0.148    0     1     1
## commute        122 104.721   68.544    1    85   300
## high_educ       135  0.415    0.495    0     0     1
## gender_num     135  0.496    0.502    0     0     1
## married_num    135  0.200    0.401    0     0     1
## high_educ_num  135  0.415    0.495    0     0     1
## children_num   135  0.148    0.357    0     0     1
## -----
```

Clean attitude

```
glimpse(attitude)
```

```
## Rows: 2,379
## Columns: 5
## $ personid   <dbl> 4122, 4122, 4122, 4122, 4122, 4122, 4122, 4122, 4122, 4122, ~
## $ year_week  <chr> "202249", "202250", "202251", "202252", "202253", "202302", ~
## $ exhaustion <dbl> 9, 8, 8, 6, 12, 12, 12, 12, 10, 12, 11, 12, 12, 12, 14, 12, ~
## $ negative   <dbl> 20, 21, 20, 17, 19, 18, 20, 16, 18, 24, 18, 19, 20, 19, 18, ~
## $ positive   <dbl> 20, 25, 24, 22, 19, 19, 19, 22, 20, 23, 23, 22, 22, 24, 23, ~
```

```
stargazer(as.data.frame(attitude), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean      St. Dev.   Min  Median  Max
## -----
## personid      2,379 33,005.050 10,331.780 4,122 36,288 45,442
## exhaustion     2,379   8.612      7.794    0.000  7.000 36.000
## negative       2,379  16.651      6.848    8.000 16.000 40.000
## positive       2,379  24.215      6.668    8.000 24.000 40.000
## -----
```

```
# checking unique values to decide int vs dbl
unique(attitude$exhaustion)
```

```
## [1] 9.00 8.00 6.00 12.00 10.00 11.00 14.00 13.00 18.00 15.00 16.00 19.00
## [13] 20.00 23.00 24.00 29.00 2.00 3.00 4.00 8.58 0.00 1.00 7.00 11.64
## [25] 6.32 5.00 6.40 17.00 25.00 27.00 7.55 22.00 9.55 21.00 31.00 9.56
```

```
## [37] 36.00 26.00 32.00 28.00 12.20 33.00 30.00 12.47 9.78 6.27 6.21 34.00
## [49] 35.00 7.04 6.86 8.55 5.99
```

```
unique(attitude$positive)
```

```
## [1] 20.00000 25.00000 24.00000 22.00000 19.00000 23.00000 21.00000 18.00000
## [9] 17.00000 8.00000 16.00000 32.00000 34.00000 30.00000 29.00000 31.00000
## [17] 24.24000 28.00000 22.56000 27.00000 26.00000 28.13000 33.00000 35.00000
## [25] 9.00000 12.00000 14.00000 13.00000 15.00000 27.19000 11.00000 10.00000
## [33] 38.00000 26.28000 36.00000 23.79000 23.48000 37.00000 39.00000 40.00000
## [41] 21.91000 22.13000 23.46000 27.18000 26.88000 27.11000 26.04000 25.04065
## [49] 27.05000
```

```
unique(attitude$negative)
```

```
## [1] 20.00000 21.00000 17.00000 19.00000 18.00000 16.00000 24.00000 28.00000
## [9] 22.00000 26.00000 25.00000 27.00000 10.00000 15.00000 14.00000 13.00000
## [17] 17.55000 11.00000 12.00000 18.53000 9.00000 17.87716 8.00000 14.85000
## [25] 15.47000 23.00000 37.00000 16.22000 30.00000 17.98000 18.19000 32.00000
## [33] 29.00000 40.00000 18.77000 31.00000 19.75000 17.57000 14.89000 15.04000
## [41] 34.00000 39.00000 36.00000 35.00000 15.55000 15.94000 16.75610 15.10000
```

We see that we do not have missing values, and the data types for all columns already fit. We only need to convert personid to a string.

```
# convert to integer to remove leading zeros and then to str
attitude$personid <- as.character(as.integer(attitude$personid))
# remove duplicates on same personid and year_week
attitude <- attitude %>%
  distinct(personid, year_week, .keep_all = TRUE)
```

Clean endperiod

```
glimpse(endperiod)
```

```
## Rows: 135
## Columns: 4
## $ personid      <dbl> 4122, 6278, 7720, 8834, 8854, 10098, 10356, 12426, 1297~
## $ promote_switch <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0~
## $ quitjob       <dbl> 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0~
## $ costofcommute <chr> "18", "12", "9", "0", "4", "12", "0", "0", "0", "0", "6.00 "~
```

```
stargazer(as.data.frame(endperiod), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean      St. Dev.   Min  Median  Max
## -----
## personid      135 32,716.780 10,835.700 4,122 37,292 45,442
## promote_switch 135  0.163      0.371      0      0      1
## quitjob       135  0.296      0.458      0      0      1
## -----
```

```
# check character columns with numerical values to extract for non numeric chars to
# be aware for special cases
```

```
find_non_numeric_chars(endperiod$costofcommute)
```

```
## [1] ". ." ". / " ". " " "NA" "CNY ."
```

```
## [7] " "      ", "      ". "      "¥."      " "
```

```
find_non_numeric_rows(endperiod$costofcommute)
```

```
## [1] "6.00 "      "- "      "25.00 / "      "11.8181819915771"
## [5] " 14.00"      "NA"      "CNY 3.00"      " 6"
## [9] "2,00"      "17.7272720336914" " 0"      "4.54545450210571"
## [13] "2,00"      "0.00 "      "¥10.00"      " 10"
## [17] "0,00"
```

here we only need to take care of the decimal dots (will be added to the extraction function), any ot

No missing values, promote_switch and quitjob are logical values, though we also will create a numerical version (0-1) for them to compute statistics. costofcommute need to be converted to doubles.

```
# convert to integer to remove leading zeros an than to str
endperiod$personid <- as.character(as.integer(endperiod$personid))
# convert to integers
endperiod$promote_switch_num <- sapply(endperiod$promote_switch, as_integer)
endperiod$quitjob_num <- sapply(endperiod$quitjob, as_integer)
# convert to logical (too have numerical and categorical variables)
endperiod$promote_switch <- as.logical(endperiod$promote_switch_num)
endperiod$quitjob <- as.logical(endperiod$quitjob_num)
# extract numbers from string. Both, dots and commas are used as decimal separators
endperiod$costofcommute<- sapply(endperiod$costofcommute, extract_numeric)
# remove duplicates on same personid
endperiod <- endperiod %>%
  distinct(personid, .keep_all = TRUE)
```

```
glimpse(endperiod)
```

```
## Rows: 135
## Columns: 6
## $ personid      <chr> "4122", "6278", "7720", "8834", "8854", "10098", "1~
## $ promote_switch <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FA~
## $ quitjob        <lgl> FALSE, TRUE, FALSE, FALSE, FALSE, TRUE, FALSE, FALS~
## $ costofcommute  <dbl> 18.00000, 12.00000, 9.00000, 0.00000, 4.00000, 12.0~
## $ promote_switch_num <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, ~
## $ quitjob_num     <int> 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, ~
```

```
stargazer(as.data.frame(endperiod), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N Mean St. Dev. Min Median Max
## -----
## promote_switch 135 0.163 0.371    0    0    1
## quitjob        135 0.296 0.458    0    0    1
## costofcommute  127 7.308 7.195  0.000 6.000 55.000
## promote_switch_num 135 0.163 0.371    0    0    1
## quitjob_num     135 0.296 0.458    0    0    1
## -----
```

Cost of commute has two missing values, which needs to be considered during analysis.

Clean performance

```
glimpse(performance)
```

```
## Rows: 9,870
## Columns: 12
## $ personid      <dbl> 4122, 4122, 4122, 4122, 4122, 4122, 4122, 4122, 4122~
## $ year_week     <chr> "202201", "202202", "202203", "202204", "202205", "2~
## $ perform1      <chr> "-1.14160859584808", "0.414592355489731", "1.0139772~
## $ phonecall     <chr> "-1.13219404220581", "0.312170952558517", "1.0971519~
## $ phonecallraw  <chr> "223", "499", "649", "709", "403", "760", "268", "73~
## $ homethatweek  <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ logphonecall  <chr> "5.40717172622681", "6.21260595321655", "6.475432872~
## $ logcallpersec <chr> "-5.10147857666016", "-5.15978050231934", "-5.142978~
## $ logcalllength <chr> "10.5086498260498", "11.372386932373", "11.618411064~
## $ logcall_dayworked <chr> "9.41003799438477", "9.58062744140625", "9.826651573~
## $ logdaysworked <chr> "1.0986123085022", "1.7917594909668", "1.79175949096~
## $ date          <chr> "2022-01-03", "2022-01-10", "2022-01-17", "2022-01-2~
```

```
stargazer(as.data.frame(performance), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean      St. Dev.  Min  Median  Max
## -----
## personid      9,870 32,493.740 10,623.780 4,122 36,908 45,442
## homethatweek  9,870   0.196     0.397      0      0      1
## -----
```

```
unique(performance$homethatweek)
```

```
## [1] 0 1
```

```
# check character columns with numerical values to extract for non numeric chars to
# be aware for special cases
```

```
find_non_numeric_chars(performance$perform1)
```

```
## [1] "." "-." ".e"
```

```
#find_non_numeric_rows(performance$perform1)
find_non_numeric_chars(performance$phonecall)
```

```
## [1] "." "-."
```

```
find_non_numeric_chars(performance$phonecallraw)
```

```
## [1] "NA" " " " " "
```

```
find_non_numeric_chars(performance$homethatweek)
```

```
## character(0)
```

```
find_non_numeric_chars(performance$logphonecall)
```

```
## [1] "." "NA" ", "
```

```
find_non_numeric_chars(performance$logcallpersec)
```

```
## [1] "." "NA" ", "
```

```
find_non_numeric_chars(performance$logcalllength)
```

```
## [1] "." "NA" ",,"
```

```
find_non_numeric_chars(performance$logcall_dayworked)
```

```
## [1] "." "NA" ",,"
```

```
find_non_numeric_chars(performance$logdaysworked)
```

```
## [1] "." ",,"
```

here we need to take care of the decimal dots, dashes to negative(negative) and scientifics notion w

Personid needs to be converted to a string, as well as all the call and performance measurements need to be converted to doubles. Homethatweek needs to be converted to a logical and numerical

```
# convert to integer to remove leading zeros an than to str
performance$personid <- as.character(as.integer(performance$personid))
# trim year_week and remove all not numerical characters
performance$year_week <- trimws(performance$year_week)
performance$year_week <- gsub("[^0-9.]", "", performance$year_week)
# convert to float/doubles
performance$perform1 <- sapply(performance$perform1, as_numeric)
performance$phonecall <- sapply(performance$phonecall, as_numeric)
# convert to integer
performance$phonecallraw <- sapply(performance$phonecallraw, as_integer)
# convert to boolean
performance$homethatweek_num <- sapply(performance$homethatweek, as_integer)
performance$homethatweek <- sapply(performance$homethatweek, as_integer)
# convert to float/doubles
performance$logphonecall <- sapply(performance$logphonecall, as_numeric)
performance$logcallpersec <- sapply(performance$logcallpersec, as_numeric)
performance$logcalllength <- sapply(performance$logcalllength, as_numeric)
performance$logcall_dayworked <- sapply(performance$logcall_dayworked, as_numeric)
performance$logdaysworked <- sapply(performance$logdaysworked, as_numeric)
# convert to Date
performance$date <- as.Date(performance$date)

# remove duplicates on same personid, year_week
performance <- performance %>%
  distinct(personid, year_week, .keep_all = TRUE)
```

```
glimpse(performance)
```

```
## Rows: 9,870
## Columns: 13
## $ personid      <chr> "4122", "4122", "4122", "4122", "4122", "4122", "412~
## $ year_week     <chr> "202201", "202202", "202203", "202204", "202205", "2~
## $ perform1      <dbl> -1.1416086, 0.4145924, 1.0139773, 1.4812315, -0.1150~
## $ phonecall      <dbl> -1.13219404, 0.31217095, 1.09715199, 1.41114438, -0.~
## $ phonecallraw   <dbl> 223, 499, 649, 709, 403, 760, 268, 73, 301, 629, 529~
## $ homethatweek   <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ logphonecall    <dbl> 5.407172, 6.212606, 6.475433, 6.563856, 5.998937, 6.~
## $ logcallpersec   <dbl> -5.101479, -5.159781, -5.142978, -5.161007, -5.15217~
## $ logcalllength   <dbl> 10.508650, 11.372387, 11.618411, 11.724862, 11.15110~
## $ logcall_dayworked <dbl> 9.410038, 9.580627, 9.826652, 9.778953, 9.359349, 9.~
```

```
## $ logdaysworked <dbl> 1.098612, 1.791759, 1.791759, 1.945910, 1.791759, 1.~
## $ date <date> 2022-01-03, 2022-01-10, 2022-01-17, 2022-01-24, 202~
## $ homethatweek_num <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
stargazer(as.data.frame(performance), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.  Min    Median  Max
## -----
## perform1      9,827 -0.015   0.988   -3.031  0.049   4.163
## phonecall     9,716 -0.006   0.963   -3.113  0.070   5.828
## phonecallraw  9,561 440.258 142.558    1     445   1,264
## homethatweek   9,870  0.196   0.397    0      0      1
## logphonecall   9,574  6.009   0.481    0.000  6.098   7.142
## logcallpersec  9,576 -5.170   0.160   -5.951 -5.169  -1.099
## logcalllength  9,576 11.175   0.521    2.485 11.273 12.116
## logcall_dayworked 9,576  9.474   0.457    2.485  9.552 10.359
## logdaysworked  9,855  1.685   0.282    0.000  1.792  1.946
## homethatweek_num 9,870  0.196   0.397    0      0      1
## -----
```

Almost all measurements have missing data.

Clean wage

```
glimpse(wage)
```

```
## Rows: 3,007
## Columns: 5
## $ basewage <chr> "1650", "1650", "1650", "1650", "1650", "1650", "1800", "18~
## $ grosswage <chr> "3650.76000976562", "4775.52001953125", "4358", "4801", "40~
## $ personid <chr> "4122", "4122", "4122", "4122", "4122", "4122", "4122", "41~
## $ wage_month <chr> "202201", "202202", "202203", "202204", "202205", "202206",~
## $ bonustotal <chr> "2000.76000976562", "3125.52001953125", "2708", "3151", "23~
stargazer(as.data.frame(wage), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic N Mean St. Dev. Min Median Max
## =====
```

```
# check character columns with numerical values to extract for non numeric chars to
# be aware for special cases
find_non_numeric_chars(wage$basewage)
```

```
## [1] ", " " " " " " " " " " " " " " " " " " " " " " " " " " " " " " " " " " "
```

```
find_non_numeric_rows(wage$basewage)
```

```
## [1] "1600,00" " 1600" " 1700" " 1600" " 1700" " 1600"
## [7] " 1600" "1800,00" "2 000.00" "1700,00" " 1700" " 1700"
## [13] " 1700" " 1700" " 1750" " 2100" "2,300.00" "1,500.00"
## [19] " 1650" " 2300" " 1700" " 1450" " 1600" "1400,00"
## [25] " 1800" "1400,00" "1400,00" "1,550.00" " 1750" "1,550.00"
## [31] " 1500" "1 400.00" " 1550" " 1800" " 1400" " 1400"
```

```
## [37] "1550,00" " 1500"      " 1750"      " 1850"      "1 400.00" " 1600"
## [43] " 1600"      "1650,00"    "1,500.00"    " 1750"      " 1500"      "1800,00"
## [49] "1 500.00" " 1500"      "1500,00"    "1,500.00"
```

```
# only trimming needed
```

We have no numerical values yet. basewage, can be converted to integer, gross_wage and bonustotal can be converted to double

```
# convert to integer to remove leading zeros an than to str
```

```
wage$personid <- as.character(as.integer(wage$personid))
wage$basewage <- sapply(wage$basewage, as_integer)
wage$grosswage <- sapply(wage$grosswage, as_integer)
wage$wage_month <- trimws(wage$wage_month)
wage$wage_month<- gsub("[^0-9.]", "", wage$wage_month)
wage <- rename(wage, year_month=wage_month)
wage$bonustotal <- sapply(wage$bonustotal, as_numeric)
# remove duplicates on same personid, yearmonth
wage <- wage %>%
  distinct(personid, year_month, .keep_all = TRUE)
```

```
glimpse(wage)
```

```
## Rows: 3,007
## Columns: 5
## $ basewage    <dbl> 1650, 1650, 1650, 1650, 1650, 1650, 1800, 1800, 1800, 1800,~
## $ grosswage   <dbl> 3650, 4775, 4358, 4801, 4045, 5497, 3123, 1890, 4725, 5918,~
## $ personid    <chr> "4122", "4122", "4122", "4122", "4122", "4122", "4122", "41~
## $ year_month  <chr> "202201", "202202", "202203", "202204", "202205", "202206",~
## $ bonustotal  <dbl> 2000.76, 3125.52, 2708.00, 3151.00, 2395.12, 3847.76, 1323.~
```

```
stargazer(as.data.frame(wage), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean      St. Dev.    Min    Median      Max
## -----
## basewage      2,955 1,666.244  221.577    650     1,600     2,850
## grosswage      2,978 3,135.297 1,066.351    48     2,979.5    14,553
## bonustotal     2,993 1,502.263   926.276    0.000  1,325.170 12,853.000
## -----
```

All columns again contain missing data. Grosswage is not the sum of basewage+bonustotal all the time:
Cases:

- NA in basewage and grosswage or basewage and bonus total or bonustotal and grosswage: remove row
- NA in basewage but grosswage and bonustotal are given: $\text{basewage} = \text{grosswage} - \text{bonustotal}$
- NA in bonustotal but grosswage and basewage are given: $\text{bonustotal} = \text{grosswage} - \text{basewage}$
- $\text{Grosswage} \neq \text{basewage} + \text{bonustotal}$: recalculate
- wage yearmonth sometimes misses the month

```
# case 2
```

```
wage <- wage %>% filter(!is.na(basewage) & !is.na(grosswage) | !is.na(bonustotal) & !is.na(grosswage) |
```

```
# case 3
```

```
wage <- wage %>% mutate(basewage = ifelse(is.na(basewage), grosswage-bonustotal, basewage))
```

```

# case 4
wage <- wage %>% mutate(bonustotal = ifelse(is.na(bonustotal), grosswage-basewage, bonustotal))

# case 5
wage <- wage %>% mutate(grosswage = ifelse(grosswage != bonustotal+basewage, bonustotal+basewage, grosswage))

# case 6
wage <- wage %>% filter(nchar(year_month)==6)

```

Merging data together

Merging

```

# First we merge all volunteer related data together on personid (unique for a person)
volunteer_endperiod <- merge(volunteer, endperiod, by="personid", all=TRUE)
# Second we merge all weekly based timeseries data together per personid
attitude_performance <- merge(attitude, performance, by=c("personid", "year_week"), all=TRUE)
# As wage only has a yearmonth column, we cannot directly merge it with an other table containing timeseries
# The goal now is to create one big dataframe with all volunteer information, their performance, attitude and wage
attitude_performance <- attitude_performance %>%
  mutate(
    week_date = ISOweek2date(
      paste0(substr(year_week,1,4), "-W", substr(year_week,5,6), "-1")
    ),
    year_month = format(week_date, "%Y%m") # gives e.g. "202401", "202412"
  )
# merge on year_month
attitude_performance_wage <- merge(attitude_performance, wage, by=c("personid", "year_month"), all.x=TRUE, all.y=FALSE)
# merge on personid
volunteer_endperiod_attitude_performance_wage <- merge(attitude_performance_wage, volunteer_endperiod, by="personid", all=TRUE)
# CAUTION: We are merging a weekly timeseries together with a monthly timeseries. When creating distributions, be careful!
# Because of the problem above, we create another dataframe with all employee related data
volunteer_endperiod_wage <- merge(volunteer_endperiod, wage, by="personid", all=TRUE)

volunteer_endperiod_wage <- volunteer_endperiod_wage %>% filter(is.finite(grosswage), is.finite(bonustotal))

# lets try to get a dataframe with ALL data and no NAs at all
distinct_all <- na.omit(volunteer_endperiod_attitude_performance_wage)

```

Analyzing merged dataframes

```
stargazer(as.data.frame(volunteer_endperiod), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N   Mean   St. Dev.  Min  Median  Max
## -----
## age            135 23.452   4.315     1    23     35
## tenure         135 23.385   24.608     2    13    150
## children       135  0.148   0.357     0     0     1
## bedroom        135  0.978   0.148     0     1     1
## commute        122 104.721  68.544     1    85    300
```

```
## high_educ      135  0.415  0.495    0    0    1
## gender_num     135  0.496  0.502    0    0    1
## married_num    135  0.200  0.401    0    0    1
## high_educ_num  135  0.415  0.495    0    0    1
## children_num   135  0.148  0.357    0    0    1
## promote_switch 135  0.163  0.371    0    0    1
## quitjob        135  0.296  0.458    0    0    1
## costofcommute  127  7.308  7.195  0.000 6.000 55.000
## promote_switch_num 135  0.163  0.371    0    0    1
## quitjob_num    135  0.296  0.458    0    0    1
## -----
```

looks good

```
stargazer(as.data.frame(attitude_performance), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.  Min    Median  Max
## -----
## exhaustion      2,379  8.612    7.794    0.000  7.000  36.000
## negative        2,379 16.651    6.848    8.000 16.000 40.000
## positive        2,379 24.215    6.668    8.000 24.000 40.000
## perform1        9,827 -0.015    0.988   -3.031  0.049  4.163
## phonecall       9,716 -0.006    0.963   -3.113  0.070  5.828
## phonecallraw    9,561 440.258 142.558    1     445  1,264
## homethatweek    9,870  0.196    0.397    0      0      1
## logphonecall    9,574  6.009    0.481    0.000  6.098  7.142
## logcallpersec   9,576 -5.170    0.160   -5.951 -5.169 -1.099
## logcalllength   9,576 11.175    0.521    2.485 11.273 12.116
## logcall_dayworked 9,576  9.474    0.457    2.485  9.552 10.359
## logdaysworked  9,855  1.685    0.282    0.000  1.792  1.946
## homethatweek_num 9,870  0.196    0.397    0      0      1
## -----
```

when using this dataframe we have to be cautious as attitude has way less columns than attitude n(att

```
stargazer(as.data.frame(attitude_performance_wage), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.  Min    Median  Max
## -----
## exhaustion      2,379  8.612    7.794    0.000  7.000  36.000
## negative        2,379 16.651    6.848    8.000 16.000 40.000
## positive        2,379 24.215    6.668    8.000 24.000 40.000
## perform1        9,827 -0.015    0.988   -3.031  0.049  4.163
## phonecall       9,716 -0.006    0.963   -3.113  0.070  5.828
## phonecallraw    9,561 440.258 142.558    1     445  1,264
## homethatweek    9,870  0.196    0.397    0      0      1
## logphonecall    9,574  6.009    0.481    0.000  6.098  7.142
## logcallpersec   9,576 -5.170    0.160   -5.951 -5.169 -1.099
## logcalllength   9,576 11.175    0.521    2.485 11.273 12.116
## logcall_dayworked 9,576  9.474    0.457    2.485  9.552 10.359
## logdaysworked  9,855  1.685    0.282    0.000  1.792  1.946
## homethatweek_num 9,870  0.196    0.397    0      0      1
## basewage        9,966 1,601.860 176.944  527.070 1,600.000 2,450.000
```

```
## grosswage      9,919 3,097.610 990.834  1,140.000 2,877.000 14,553.000
## bonustotal     9,966 1,496.445 913.240    0.000   1,285.725 12,853.000
## -----
```

```
stargazer(as.data.frame(volunteer_endperiod_attitude_performance_wage), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## exhaustion      2,379    8.612    7.794      0.000     7.000     36.000
## negative        2,379   16.651    6.848      8.000    16.000     40.000
## positive        2,379   24.215    6.668      8.000    24.000     40.000
## perform1        9,827   -0.015    0.988     -3.031     0.049     4.163
## phonecall       9,716   -0.006    0.963     -3.113     0.070     5.828
## phonecallraw    9,561   440.258  142.558      1       445     1,264
## homethatweek    9,870    0.196    0.397      0         0         1
## logphonecall    9,574    6.009    0.481      0.000     6.098     7.142
## logcallpersec   9,576   -5.170    0.160     -5.951    -5.169    -1.099
## logcalllength   9,576   11.175    0.521      2.485    11.273    12.116
## logcall_dayworked 9,576    9.474    0.457      2.485     9.552    10.359
## logdaysworked   9,855    1.685    0.282      0.000     1.792     1.946
## homethatweek_num 9,870    0.196    0.397      0         0         1
## basewage        9,966  1,601.860 176.944   527.070  1,600.000 2,450.000
## grosswage       9,919  3,097.610 990.834  1,140.000 2,877.000 14,553.000
## bonustotal      9,966  1,496.445 913.240    0.000   1,285.725 12,853.000
## age             10,079   23.485    4.191      1        23        35
## tenure          10,079   23.893   24.773      2        18       150
## children         10,079    0.141    0.348      0         0         1
## bedroom          10,079    0.974    0.160      0         1         1
## commute          9,118  103.519   67.620      1        80       300
## high_educ        10,079    0.400    0.490      0         0         1
## gender_num       10,079    0.498    0.500      0         0         1
## married_num      10,079    0.181    0.385      0         0         1
## high_educ_num    10,079    0.400    0.490      0         0         1
## children_num     10,079    0.141    0.348      0         0         1
## promote_switch   10,079    0.182    0.386      0         0         1
## quitjob          10,079    0.238    0.426      0         0         1
## costofcommute    9,488    7.213    7.231      0.000     6.000    55.000
## promote_switch_num 10,079    0.182    0.386      0         0         1
## quitjob_num      10,079    0.238    0.426      0         0         1
## -----
```

```
stargazer(as.data.frame(volunteer_endperiod_wage), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## age            2,968   23.429    4.476      1        23        35
## tenure         2,968   24.113   24.607      2        19       150
## children       2,968    0.141    0.348      0         0         1
## bedroom        2,968    0.975    0.155      0         1         1
## commute        2,698  103.935   67.178      1        80       300
## high_educ      2,968    0.412    0.492      0         0         1
```

```
## gender_num      2,968  0.510    0.500    0      1      1
## married_num     2,968  0.190    0.392    0      0      1
## high_educ_num   2,968  0.412    0.492    0      0      1
## children_num    2,968  0.141    0.348    0      0      1
## promote_switch  2,968  0.190    0.392    0      0      1
## quitjob         2,968  0.202    0.402    0      0      1
## costofcommute   2,770  7.097    6.983    0.000    6.000    55.000
## promote_switch_num 2,968  0.190    0.392    0      0      1
## quitjob_num     2,968  0.202    0.402    0      0      1
## basewage        2,968 1,665.141 224.524 248.000 1,600.000 2,850.000
## grosswage       2,968 3,169.689 1,025.056 53.000 2,986.500 14,553.000
## bonustotal      2,968 1,504.548 931.247 -1,597.000 1,328.080 12,853.000
## -----
```

```
stargazer(as.data.frame(distinct_all), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## exhaustion     1,785   8.983    8.105    0.000    8.000    36.000
## negative       1,785  16.699    7.088    8.000   16.000    40.000
## positive       1,785  23.718    6.730    8.000   24.000    40.000
## perform1       1,785   0.070    0.866   -2.906    0.096    2.772
## phonecall      1,785   0.030    0.814   -3.098    0.046    2.266
## phonecallraw   1,785 439.319 126.850     2      442     835
## homethatweek   1,785   0.508    0.500     0       1       1
## logphonecall   1,785   6.032    0.367    0.693    6.091    6.727
## logcallpersec  1,785  -5.158    0.129   -5.541   -5.157   -4.075
## logcalllength  1,785  11.191    0.375    5.561   11.242   11.952
## logcall_dayworked 1,785   9.465    0.354    3.951    9.506   10.104
## logdaysworked 1,785   1.726    0.182    0.693    1.792    1.946
## homethatweek_num 1,785   0.508    0.500     0       1       1
## basewage       1,785 1,683.305 158.587 1,300.000 1,650.000 2,150.000
## grosswage      1,785 3,327.999 1,046.363 1,740.000 3,095.000 14,553.000
## bonustotal     1,785 1,644.695 969.980  40.000   1,408.500 12,853.000
## age           1,785  23.519    3.254    18      23      34
## tenure        1,785  23.685   26.683     2      13     150
## children       1,785   0.076    0.264     0       0       1
## bedroom        1,785   0.941    0.235     0       1       1
## commute        1,785 104.871   67.073    20      80     300
## high_educ      1,785   0.345    0.475     0       0       1
## gender_num     1,785   0.459    0.498     0       0       1
## married_num    1,785   0.111    0.314     0       0       1
## high_educ_num  1,785   0.345    0.475     0       0       1
## children_num   1,785   0.076    0.264     0       0       1
## promote_switch  1,785   0.182    0.386     0       0       1
## quitjob        1,785   0.000    0.000     0       0       0
## costofcommute  1,785   6.749    6.160    0.000    5.000   30.000
## promote_switch_num 1,785   0.182    0.386     0       0       1
## quitjob_num    1,785   0.000    0.000     0       0       0
## -----
```

When joining everything together we can see that we do not have all data for each person and its work weeks, what was to be expected. When removing all NAs we can see that the numbers slightly change and for exam-

ples all quitters are gone. Therefore we have to use the `volunteer_endperiod_attitude_performance_wage` with caution and eliminate NAs for each calculation individually. Though this dataframe is handy to use for the further work.

`volunteer_enperiod`: Cross sectional data of all employee related information.

`attitude_performance`: Panel data on how an employee performed in a certain week of the year and how he/she felt.

`attitude_performance_wage`: Panel data on how an employee performed in a certain week of the year, how he/she felt in that week and their wage during that month. Caution: contains two different time series.

`volunteer_endperiod_attitude_performance_wage`: Like `attitude_performance_wage`, but with `volunteer_endperiod` data.

General Analysis

Population: All employees of Huaxia

Our sample: Depending on the table we use and what we collect the unit of interest might differ:

Attitude, performance (time series): weekly observations of a person

volunteer (cross sectional): the person itself

wage (time series): the wage of a person in a month

endperiod (cross sectional): behavior of a person(?)

`volunteer_endperiod_attitude_performance_wage`(panel data): repeated observations of different employees across multiple time periods, combining attitudes performance, volunteer, endperiod, and wages

Get first basic insights of the tables before eliminating null values and merging data together. For a better overview we will look at each table alone, not on the composition table.

Statistics

Performance

```
stargazer(as.data.frame(performance), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.   Min    Median   Max
## -----
## perform1       9,827  -0.015    0.988   -3.031  0.049   4.163
## phonecall      9,716  -0.006    0.963   -3.113  0.070   5.828
## phonecallraw   9,561  440.258  142.558     1     445   1,264
## homethatweek   9,870   0.196    0.397     0      0      1
## logphonecall   9,574   6.009    0.481    0.000  6.098   7.142
## logcallpersec  9,576  -5.170    0.160   -5.951 -5.169  -1.099
## logcalllength  9,576  11.175    0.521    2.485  11.273  12.116
## logcall_dayworked 9,576   9.474    0.457    2.485  9.552  10.359
## logdaysworked 9,855   1.685    0.282    0.000  1.792   1.946
## homethatweek_num 9,870   0.196    0.397     0      0      1
## -----
```

```
df<-performance %>% filter(is.finite(perform1))
# perform plot
stats <- df %>% summarise(
  min  = min(perform1, na.rm = TRUE),
  max  = max(perform1, na.rm = TRUE),
  mean = mean(perform1, na.rm = TRUE),
```

```

median = median(perform1, na.rm = TRUE),
sd      = sd(perform1, na.rm = TRUE),
n       = n()
)

p<-ggplot(df, aes(y = perform1)) +
  geom_boxplot(
    outlier.alpha = 0.6,
    outlier.size = 1.8,
    fill = "#9df3c4",
    color = "#0d0c3f"
  ) +
  labs(y = "Performance Score", title = "Distribution of performances per week of employees") +
  theme_minimal() +
  theme(axis.title.x = element_blank(),
        axis.text.x = element_blank(),
        axis.ticks.x = element_blank())


# add annotations
p + annotate("text", x = 0.27, y = 4,
             label = paste0("SD: ", round(stats$sd,2)),
             hjust = 0, color = "#0d0c3f")+
  annotate("text", x = 0.27, y = 3.6,
             label = paste0("Mean: ", round(stats$mean,2)),
             text = 0, color = "#0d0c3f")+
  annotate("text", x = 0.27, y = 3.2,
             label = paste0("Median: ", round(stats$median,2)),
             hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 0.27, y = 2.8,
             label = paste0("Min: ", round(stats$min,2)),
             hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 0.27, y = 2.4,
             label = paste0("Max: ", round(stats$max,2)),
             hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 0.27, y = 2.0,
             label = paste0("N: ", round(stats$n,2)),
             hjust = 0, color = "#0d0c3f")

```

```

## Warning in annotate("text", x = 0.27, y = 3.6, label = paste0("Mean: ", :
## Ignoring unknown parameters: `text`

```



```
ggsave("performance_distribution.png", width = 7, height = 5, dpi = 300)
```

A lot of missing values, 296 at max, which will fall out from the data set when working with it.
Working from home rate is at 19,64% on average, probably as a result of a Work-from-home-policy.

Volunteer

```
stargazer(as.data.frame(volunteer), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N   Mean   St. Dev. Min Median Max
## -----
## age            135 23.452   4.315    1    23   35
## tenure         135 23.385  24.608    2    13  150
## children       135  0.148   0.357    0     0    1
## bedroom        135  0.978   0.148    0     1    1
## commute        122 104.721  68.544    1    85  300
## high_educ      135  0.415   0.495    0     0    1
## gender_num     135  0.496   0.502    0     0    1
## married_num    135  0.200   0.401    0     0    1
## high_educ_num  135  0.415   0.495    0     0    1
## children_num   135  0.148   0.357    0     0    1
## -----
```

No missing values. The average worker is 23,5 years old with a mean commute time of 100 minutes per week while there are greater differences, and is for almost 23 months within the company. 15% have children, 20% are married, more than 40% have a higher education. Both, men and women work at the company.

Endperiod

```
stargazer(as.data.frame(endperiod), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N  Mean  St. Dev.  Min  Median  Max
## -----
## promote_switch  135 0.163  0.371    0    0      1
## quitjob         135 0.296  0.458    0    0      1
## costofcommute   127 7.308  7.195   0.000 6.000  55.000
## promote_switch_num 135 0.163  0.371    0    0      1
## quitjob_num     135 0.296  0.458    0    0      1
## -----
```

On average, an employee had cost of commute of 7.5 but this number may heavily differ to some employees, 16% got a promotion, while almost 30% quit their job at some point.

Wage

```
stargazer(as.data.frame(wage), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## basewage      2,981 1,665.260  224.347   248.000   1,600.000 2,850.000
## grosswage     2,968 3,169.689 1,025.056   53.000   2,986.500 14,553.000
## bonustotal    2,981 1,504.046   930.508 -1,597.000 1,329.380 12,853.000
## -----
```

```
df<-wage %>% filter(is.finite(grosswage))
stats <- df %>% filter(is.finite(grosswage))%>% summarise(
  min = min(grosswage, na.rm = TRUE),
  max = max(grosswage, na.rm = TRUE),
  mean = mean(grosswage, na.rm = TRUE),
  median = median(grosswage, na.rm = TRUE),
  sd = sd(grosswage, na.rm = TRUE),
  n = n()
)

p <- ggplot(df, aes(y = grosswage)) +
  geom_boxplot(
    outlier.alpha = 0.6,
    outlier.size = 1.8,
    fill = "#9df3c4",
    color = "#0d0c3f"
  ) +
  labs(y = "€ per month", title = "Distribution of Gross Wages") +
  theme_minimal()

# add annotations
p + annotate("text", x = 0.27, y = 14000,
  label = paste0("SD: ", round(stats$sd,0), "€"),
  hjust = 0, color = "#0d0c3f")+
```

```

annotate("text", x = 0.27, y = 13000,
        label = paste0("Mean: ", round(stats$mean,0), "€"),
        hjust = 0, color = "#0d0c3f") +
annotate("text", x = 0.27, y = 12000,
        label = paste0("Median: ", round(stats$median,0), "€"),
        hjust = 0, color = "#0d0c3f") +
annotate("text", x = 0.27, y = 11000,
        label = paste0("Min: ", round(stats$min,0), "€"),
        hjust = 0, color = "#0d0c3f") +
annotate("text", x = 0.27, y = 10000,
        label = paste0("Max: ", round(stats$max,0), "€"),
        hjust = 0, color = "#0d0c3f") +
annotate("text", x = 0.27, y = 9000,
        label = paste0("N: ", round(stats$n,2)),
        hjust = 0, color = "#0d0c3f")

```



```

ggsave("gross_wage_distribution.png", width = 7, height = 5, dpi = 300)

```

On average an employee had a base wage of around 1650€ which can slightly differ. A heavily uneven distribution can be seen for the bonuses for employees, while some refer no bonus at all others receive almost 13.000€ as bonus, on average it is 900€ per employee.

Attitude

```

stargazer(as.data.frame(attitude), type = "text", median = TRUE, header = TRUE)

```

```
##
```

```
## =====
## Statistic      N      Mean  St. Dev.  Min  Median  Max
## -----
## exhaustion 2,379 8.612   7.794   0.000 7.000 36.000
## negative   2,379 16.651  6.848   8.000 16.000 40.000
## positive   2,379 24.215  6.668   8.000 24.000 40.000
## -----
```

Exhaustion scores are mostly low and clustered around 7–9 but some people reach much higher (up to 36), negative scores center near 16 with moderate variation across the group, and positive scores center near 24 with a similar moderate spread, meaning most participants stay within roughly one standard deviation (about ± 7 points) of these averages but a few reach the extreme minimum or maximum values

All together

```
stargazer(as.data.frame(volunteer_endperiod_attitude_performance_wage), type = "text", median = TRUE, h
```

```
##
## =====
## Statistic      N      Mean  St. Dev.  Min  Median  Max
## -----
## exhaustion      2,379   8.612   7.794   0.000   7.000  36.000
## negative        2,379  16.651   6.848   8.000  16.000  40.000
## positive        2,379  24.215   6.668   8.000  24.000  40.000
## perform1        9,827  -0.015   0.988  -3.031   0.049   4.163
## phonecall       9,716  -0.006   0.963  -3.113   0.070   5.828
## phonecallraw    9,561  440.258  142.558    1     445   1,264
## homethatweek    9,870   0.196   0.397    0      0      1
## logphonecall    9,574   6.009   0.481   0.000   6.098   7.142
## logcallpersec   9,576  -5.170   0.160  -5.951  -5.169  -1.099
## logcalllength   9,576  11.175   0.521   2.485  11.273  12.116
## logcall_dayworked 9,576   9.474   0.457   2.485   9.552  10.359
## logdaysworked  9,855   1.685   0.282   0.000   1.792   1.946
## homethatweek_num 9,870   0.196   0.397    0      0      1
## basewage        9,966  1,601.860  176.944  527.070 1,600.000 2,450.000
## grosswage       9,919  3,097.610  990.834 1,140.000 2,877.000 14,553.000
## bonustotal      9,966  1,496.445  913.240   0.000  1,285.725 12,853.000
## age            10,079  23.485   4.191    1     23     35
## tenure          10,079  23.893   24.773    2     18    150
## children        10,079   0.141   0.348    0      0      1
## bedroom         10,079   0.974   0.160    0      1      1
## commute         9,118  103.519  67.620    1     80    300
## high_educ       10,079   0.400   0.490    0      0      1
## gender_num      10,079   0.498   0.500    0      0      1
## married_num     10,079   0.181   0.385    0      0      1
## high_educ_num   10,079   0.400   0.490    0      0      1
## children_num    10,079   0.141   0.348    0      0      1
## promote_switch  10,079   0.182   0.386    0      0      1
## quitjob         10,079   0.238   0.426    0      0      1
## costofcommute   9,488   7.213   7.231   0.000   6.000  55.000
## promote_switch_num 10,079   0.182   0.386    0      0      1
## quitjob_num     10,079   0.238   0.426    0      0      1
## -----
```

```

df <- volunteer_endperiod_attitude_performance_wage
df <- dplyr::select(df, where(is.numeric))

# create a summary df as preparation for the radar plot rendering
create_summary <- function(df, drop_all_na = TRUE, keep_group_index = TRUE) {
  num <- df %>%
    select(where(is.numeric)) %>%
    mutate(across(everything(), ~ ifelse(is.finite(.x), .x, NA_real_)))

  if (drop_all_na) {
    keep <- vapply(num, function(x) any(!is.na(x)), logical(1))
    num <- num[, keep, drop = FALSE]
  }
  stopifnot(ncol(num) > 0)

  summary <- num %>%
    summarise(
      across(everything(),
        list(mean = ~mean(.x, na.rm = TRUE),
              min = ~min(.x, na.rm = TRUE),
              max = ~max(.x, na.rm = TRUE)),
        .names = "{.col}_{.fn}")
    )
  summary
}

# Expects a dataframe with *_min, *_max, *_mean columns
render_radar_plot <- function (df, pad = 1e-6) {
  # separate min, max, mean
  max_vals <- df %>% select(ends_with("_max")) %>% rename_with(~ sub("_max$", "", .x))
  min_vals <- df %>% select(ends_with("_min")) %>% rename_with(~ sub("_min$", "", .x))
  mean_vals <- df %>% select(ends_with("_mean")) %>% rename_with(~ sub("_mean$", "", .x))

  # avoid zero-range axes (min == max) that break fmsb scaling
  rng <- as.numeric(max_vals[1, ] - min_vals[1, ])
  z <- which(is.na(rng) | rng == 0)
  if (length(z)) {
    max_vals[1, z] <- as.numeric(max_vals[1, z]) + pad
    min_vals[1, z] <- as.numeric(min_vals[1, z]) - pad
  }

  web <- rbind(max_vals, min_vals, mean_vals)
  web[] <- lapply(web, as.numeric)

  radarchart(web, na.itp = TRUE)
}

radar_by_group <- function(df, group_col = "group_index", stat = "mean") {
  num <- df %>%
    mutate(across(where(is.numeric), ~ ifelse(is.finite(.x), .x, NA_real_)))

  # 1) aggregate per group

```

```

fun <- switch(stat,
  mean = ~mean(.x, na.rm = TRUE),
  median = ~median(.x, na.rm = TRUE),
  stop("stat must be 'mean' or 'median'"))
grp <- num %>%
  group_by(.data[[group_col]]) %>%
  summarise(across(where(is.numeric), fun), .groups = "drop")

# 2) build fmsb input: first two rows = per-variable max/min over all groups
vals <- grp %>% select(-all_of(group_col))
rownames(vals) <- grp[[group_col]]
max_vals <- summarise_all(vals, max, na.rm = TRUE)
min_vals <- summarise_all(vals, min, na.rm = TRUE)

# pad zero-range axes
rng <- as.numeric(max_vals[1, ] - min_vals[1, ])
z <- which(is.na(rng) | rng == 0)
if (length(z)) {
  max_vals[1, z] <- as.numeric(max_vals[1, z]) + 1e-6
  min_vals[1, z] <- as.numeric(min_vals[1, z]) - 1e-6
}

web <- rbind(max_vals, min_vals, vals)


# 3) colors and legend
n <- nrow(web) - 2
pal <- brewer.pal(max(3, min(8, n)), "Set2")[seq_len(n)]
fill <- adjustcolor(pal, alpha.f = 0.25)

radarchart(
  web,
  pcol = pal,
  pfcol = fill,
  plwd = 2,
  plty = 1,
  cglcol = "grey80",
  cglty = 1,
  cglwd = 0.8,
  axislabcol = "grey30",
  vlce = 0.8
)

legend("topright",
  legend = rownames(web)[-c(1,2)],
  bty = "n", pch = 15, pt.cex = 1.3,
  col = fill, text.col = "grey20")
}

summary <- create_summary(df)
render_radar_plot(summary)

```





Basic correlation analysis

```
nums <- dplyr::select(volunteer_endperiod_attitude_performance_wage, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
```

```
## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is
## zero
```

```
ggcorrplot(cm, lab = FALSE)
```



Here we can see that the performance metric “perform1” has a correlation to phonecall, phonecallraw,

logphonecall, logcalllength, which gives us a hint on how this metrics was calculated. Furthermore we can see, a negative correlation between “negative” and “positive” from the attitude dataset, which makes sense as these are opposites. Because of the timeseries data of the performance data set, the light correlation can be explained as the attitude may change from day to day, or not (what we can not infer from that). There is also a slight correlation between exhaustion and negative attitude, which seems logical. What we cannot see are significant correlations between quit_job and any variable and other than the cluster itself between perform1 and any other variable. Quit_job has moderate negative correlations to bonuses/grosswage and promotions.

We cannot see any clear indicators on how to fight churn rate and/or what motivates people to perform better.

Detailed exploration analysis

Function for radar plots

```
# A function to render web plots for our data with some modificatino parameters
# The function first selects the desired column and than summarises the data based on what groups to co
radar_by_group <- function(df, include_group_indices = c(), group_col = "group_index", stat = "mean", ti
  # open device to save as png
  if(save==TRUE){
    png(paste0(title, ".png"), width = 2400, height = 2000, res = 300)
  }
  # check for parameters
  stopifnot(group_col %in% names(df))
  # select, based on parameters, columns
  final_selection = c(group_col, "gender_num", "married_num", "bedroom", "children_num", "high_educ_num"
  if (show_full_attitude==TRUE) {
    selection = c("negative", "positive")
    final_selection <- c(selection, final_selection)
  }
  if (show_full_performance==TRUE) {
    selection = c("logdaysworked", "logcall_dayworked", "logcalllength", "logcallpersec", "logphonecall", "p
    final_selection <- c(selection, final_selection)
  }
  if (show_full_wage==TRUE) {
    selection = c("bonustotal", "basewage")
    final_selection <- c(selection, final_selection)
  }
  if (show_commute==TRUE) {
    selection = c("commute", "costofcommute")
    final_selection <- c(selection, final_selection)
  }
  # chek if group col is in any of the selected field, if so, remove it
  if (group_col %in% sub("_num$", "", final_selection)){
    final_selection <- final_selection[final_selection!=paste0(group_col, "_num")]
  }
  print(final_selection)
  df <- df %>% select(final_selection)
  # only used numeric columns (+ group column) for computation and only real numbers
  num <- df %>%
    select(all_of(group_col), where(is.numeric)) %>%
    mutate(across(where(is.numeric), ~ ifelse(is.finite(.x), .x, NA_real_)))
```

```

# remove "_num" to shorten names of var
num <- num %>% rename_with(~ sub("_num$", "", .x))
# select the min max values based on the 90%/10% quantile of the overall data
max_vals <- summarise(num, across(where(is.numeric), ~ quantile(.x, 0.9, na.rm = TRUE)))
min_vals <- summarise(num, across(where(is.numeric), ~ quantile(.x, 0.1, na.rm = TRUE)))
# aggregate per group (mean or median)
agg_fun <- switch(stat,
  mean = ~mean(.x, na.rm = TRUE),
  median = ~median(.x, na.rm = TRUE),
  stop("stat must be 'mean' or 'median'"))
vals <- num %>%
  group_by(.data[[group_col]]) %>%
  summarise(across(where(is.numeric), agg_fun), .groups = "drop")

# get the correct label
if (length(include_group_indices)>0){
  print(vals)
  print(group_col)
  print(include_group_indices)
  labels <- vals %>% filter(.data[[group_col]] %in% include_group_indices)
  print(labels)
  labels <- as.character(labels[[group_col]])
  print(labels)
}
# again check for not numeric values and drop them if needed
keep <- vapply(vals, function(x) any(!is.na(x)), logical(1))
vals <- vals[, keep, drop = FALSE]
# check for enough values to display
if (ncol(vals) < 3) stop("Need 3 valid numeric variables after cleaning.")
# fill remaining NA cells with column means to avoid NA polygons
cm <- vapply(vals, function(x) mean(x, na.rm = TRUE), numeric(1))
for (j in seq_along(vals)) vals[[j]][is.na(vals[[j]])] <- cm[j]
# per-variable max/min, pad zero ranges
rng <- as.numeric(max_vals[1, ] - min_vals[1, ])
z <- which(is.na(rng) | rng == 0)
pad <- 1e-6
rng <- as.numeric(max_vals[1, ] - min_vals[1, ])
z <- which(is.na(rng) | rng == 0)
if (length(z)) {
  mx <- as.numeric(max_vals[1, ])
  mn <- as.numeric(min_vals[1, ])
  mx[z] <- mx[z] + pad
  mn[z] <- mn[z] - pad
  max_vals[1, ] <- as.list(mx)
  min_vals[1, ] <- as.list(mn)
}
# ensure base data.frames to simplify row assignment
max_vals <- as.data.frame(max_vals)
min_vals <- as.data.frame(min_vals)
# sub select values from desired groups
if (length(include_group_indices)>0) {
  vals <- vals %>% filter(.data[[group_col]] %in% include_group_indices)
}

```

```

# remove group index to not be displayed as category in the web chart
vals <- vals %>% select(-all_of(group_col))
if (group_col %in% colnames(max_vals)) {
  max_vals <- max_vals %>% select(-all_of(group_col))
}
if (group_col %in% colnames(min_vals)) {
  min_vals <- min_vals %>% select(-all_of(group_col))
}
# cap values if they exceed the bounds
vals <- vals %>%
  mutate(across(
    everything(),
    ~ pmax(
      pmin(.x, max_vals[[cur_column()]] [1]), # upper bound for this var
      min_vals[[cur_column()]] [1]           # lower bound for this var
    )
  ))
# put everything together
web <- rbind(max_vals, min_vals, vals)
web[] <- lapply(web, function(x) as.numeric(x)) # ensure numeric matrix

# # make it look nice
n <- nrow(web) - 2
cols <- grDevices::hcl(h = seq(15, 375, length.out = n + 1)[1:n],
  c = 80, l = 60)
fills <- grDevices::adjustcolor(cols, alpha.f = 0.25)
cols <- c("#2CB1A6", # teal-mint
  "#1B1E3F", # deep navy
  "#94A3B8", # cool gray
  "#A7F3D0", # light mint
  "#F4A261") # muted coral
fills <- adjustcolor(cols, alpha.f = 0.25)
# plot it
radarchart(web,
  pcol = cols[seq_len(nrow(web)-2)],
  pfc = fills[seq_len(nrow(web)-2)],
  plwd = 2, plty = 1,
  cglcol = "grey80", cglwd = 0.8, axislabcol = "grey30",
  vlce = 1.2, na.itp = TRUE)
title(main = title)
# add a little legend
if(!is.null(group_col)){
  legend("topright", legend = labels,
    bty = "n", pch = 15, pt.cex = 1.3, col = cols[seq_len(nrow(web)-2)], text.col = "grey20")
}
# close device
if(save==TRUE){
  dev.off()
}
}

```

Quitter vs no-quitters

Function to aggregate values to monthly timeseries

```
to_monthly <- function(df){
  df <- df %>% group_by(personid, year_month) %>% summarise(
    costofcommute = mean(costofcommute, na.rm = TRUE),
    positive = mean(positive, na.rm = TRUE),
    negative = mean(negative, na.rm = TRUE),
    exhaustion = mean(exhaustion, na.rm = TRUE),
    bedroom = mean(bedroom, na.rm = TRUE),
    promote_switch = mean(promote_switch, na.rm = TRUE),
    promote_switch_num = mean(promote_switch_num, na.rm = TRUE),
    quitjob_num = mean(quitjob_num, na.rm = TRUE),
    commute = mean(commute, na.rm = TRUE),
    tenure = mean(tenure, na.rm = TRUE),
    age = mean(age, na.rm = TRUE),
    grosswage = mean(grosswage, na.rm = TRUE),
    basewage = mean(basewage, na.rm = TRUE),
    bonustotal = mean(bonustotal, na.rm = TRUE),
    logdaysworked = mean(logdaysworked, na.rm = TRUE),
    logcall_dayworked = mean(logcall_dayworked, na.rm = TRUE),
    logcallpersec = mean(logcallpersec, na.rm = TRUE),
    logphonecall = mean(logphonecall, na.rm = TRUE),
    logcalllength = mean(logcalllength, na.rm = TRUE),
    phonecallraw = mean(phonecallraw, na.rm = TRUE),
    phonecall = mean(phonecall, na.rm = TRUE),
    perform1 = mean(perform1, na.rm = TRUE),
    children_num = mean(children_num, na.rm = TRUE),
    high_educ_num = mean(high_educ_num, na.rm = TRUE),
    married_num = mean(married_num, na.rm = TRUE),
    gender_num = mean(gender_num, na.rm = TRUE),
    homethatweek_num = mean(homethatweek_num, na.rm = TRUE),
    quitjob = first(quitjob),
    married = first(married),
    high_educ = first(high_educ),
    bedroom = first(bedroom),
    children = first(children),
    promote_switch = first(promote_switch),
    homethatweek = first(homethatweek),
    gender = first(gender),
  )
  df %>% ungroup()
}
```

Lets compare quitter vs non-quitter to better understand both groups:

```
df <- volunteer_endperiod_attitude_performance_wage
# use monthly, performance and attitude data to make compareable with wage
df <- to_monthly(df)
```

```
## `summarise()` has grouped output by 'personid'. You can override using the
## `.groups` argument.
```

```

quitter_df <- df %>% filter(quitjob == TRUE)
non_quitter_df <- df %>% filter(quitjob == FALSE)

# general impression
stargazer(as.data.frame(quitter_df), type = "text", median = TRUE, header = TRUE)

```

```

##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute  571    9.104    10.349    0.000     6.000    55.000
## bedroom        590     1.000     0.000     1         1         1
## promote_switch  590     0.000     0.000     0         0         0
## promote_switch_num 590     0.000     0.000     0         0         0
## quitjob_num     590     1.000     0.000     1         1         1
## commute        519   104.037    76.522     1        120        300
## tenure         590    23.244    28.449     2         10        150
## age            590    23.690     3.116    19         23         32
## grosswage      577  2,655.873  796.984  1,140.000  2,493.030 10,830.070
## basewage       579  1,560.781  189.256   527.070   1,550.000  2,450.000
## bonustotal     579  1,096.817  730.332    0.000    960.030   9,330.070
## logdaysworked  589     1.608     0.281    0.000     1.655     1.946
## logcall_dayworked 563     9.361     0.582    2.485     9.464    10.129
## logcallpersec   563    -5.140     0.233   -5.693    -5.148    -1.099
## logphonecall    565     5.880     0.565    1.386     5.985     6.834
## logcalllength   563    11.004     0.681    2.485    11.138    11.773
## phonecallraw    565   409.055   126.195    4.000   407.200   973.000
## phonecall       573    -0.271     0.892   -3.084    -0.211     3.784
## perform1       586    -0.326     0.868   -3.031    -0.272     2.391
## children_num    590     0.225     0.418     0         0         1
## high_educ_num   590     0.468     0.499     0         0         1
## married_num     590     0.242     0.429     0         0         1
## gender_num      590     0.451     0.498     0         0         1
## homethatweek_num 590     0.090     0.278    0.000    0.000     1.000
## quitjob        590     1.000     0.000     1         1         1
## high_educ      590     0.468     0.499     0         0         1
## children       590     0.225     0.418     0         0         1
## homethatweek   590     0.090     0.286     0         0         1
## -----

```

```

stargazer(as.data.frame(non_quitter_df), type = "text", median = TRUE, header = TRUE)

```

```

##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute  1,680     6.592     5.685    0.000     6.000    30.000
## positive       549    24.242     5.826    8.000    24.250    40.000
## negative       549    16.635     6.071    8.000    16.500    40.000
## exhaustion     549     8.575     7.036    0.000     8.000    36.000
## bedroom       1,807     0.967     0.179     0         1         1
## promote_switch 1,807     0.240     0.427     0         0         1
## promote_switch_num 1,807     0.240     0.427     0         0         1
## quitjob_num    1,807     0.000     0.000     0         0         0

```

```
## commute      1,650  103.427  64.554      2      80      300
## tenure       1,807  24.080  23.433      2      19      150
## age          1,807  23.418   4.534      1      23       35
## grosswage    1,779 3,213.903 1,026.384 1,264.430 2,997.000 14,553.000
## basewage     1,787 1,614.684 173.450   650.000 1,600.000 2,300.000
## bonustotal   1,787 1,599.423 948.947    0.000 1,373.590 12,853.000
## logdaysworked 1,807   1.694   0.193    0.000   1.729   1.946
## logcall_dayworked 1,790   9.479   0.385    3.563   9.542  10.124
## logcallpersec 1,790  -5.180   0.133   -5.686  -5.174  -4.169
## logphonecall  1,786   6.017   0.371    0.000   6.075   6.805
## logcalllength 1,790  11.188   0.411    5.172  11.252  11.854
## phonecallraw  1,786 443.698  111.384    1.000 447.267  907.000
## phonecall     1,796   0.020   0.750   -3.113   0.052   2.447
## perform1      1,804   0.031   0.795   -3.031   0.063   2.846
## children_num  1,807   0.117   0.322     0       0       1
## high_educ_num 1,807   0.387   0.487     0       0       1
## married_num   1,807   0.168   0.374     0       0       1
## gender_num    1,807   0.515   0.500     0       1       1
## homethatweek_num 1,807   0.235   0.417    0.000   0.000   1.000
## quitjob       1,807   0.000   0.000     0       0       0
## high_educ     1,807   0.387   0.487     0       0       1
## children      1,807   0.117   0.322     0       0       1
## homethatweek  1,753   0.228   0.420     0       0       1
## -----
```


```
stargazer(as.data.frame(df), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute  2,251   7.229    7.242    0.000    6.000    55.000
## positive       549   24.242    5.826    8.000   24.250   40.000
## negative       549   16.635    6.071    8.000   16.500   40.000
## exhaustion     549    8.575    7.036    0.000    8.000   36.000
## bedroom       2,397    0.975    0.156     0        1        1
## promote_switch 2,397    0.181    0.385     0         0        1
## promote_switch_num 2,397    0.181    0.385     0         0        1
## quitjob_num    2,397    0.246    0.431     0         0        1
## commute        2,169  103.573   67.593     1        80       300
## tenure         2,397  23.874   24.759     2        15       150
## age            2,397  23.485    4.231     1        23        35
## grosswage      2,356 3,077.238 1,004.152 1,140.000 2,860.500 14,553.000
## basewage       2,366 1,601.493 178.915   527.070 1,600.000 2,450.000
## bonustotal     2,366 1,476.427 925.801    0.000 1,271.050 12,853.000
## logdaysworked 2,396    1.673    0.221    0.000    1.701    1.946
## logcall_dayworked 2,353    9.451    0.443    2.485    9.524   10.129
## logcallpersec  2,353   -5.171    0.164   -5.693   -5.167   -1.099
## logphonecall   2,351    5.984    0.430    0.000    6.052    6.834
## logcalllength  2,353   11.144    0.495    2.485   11.222   11.854
## phonecallraw   2,351 435.372  116.038    1.000 436.250  973.000
## phonecall      2,369  -0.050    0.796   -3.113  -0.022    3.784
## perform1       2,390  -0.056    0.828   -3.031  -0.033    2.846
## children_num   2,397    0.144    0.351     0         0        1
## high_educ_num  2,397    0.407    0.491     0         0        1
```


```
## married_num      2,397    0.186    0.389    0      0      1
## gender_num       2,397    0.499    0.500    0      0      1
## homethatweek_num 2,397    0.199    0.392    0.000    0.000    1.000
## quitjob          2,397    0.246    0.431    0      0      1
## high_educ        2,397    0.407    0.491    0      0      1
## children         2,397    0.144    0.351    0      0      1
## homethatweek     2,343    0.193    0.395    0      0      1
## -----
```

```
# plot per numeric variable
```


```
df<-df %>% filter(is.finite(perform1), is.finite(commute), is.finite(costofcommute))
ggplot(df, aes(x=quitjob, y=perform1)) + geom_boxplot()
```




```
ggplot(df, aes(x=quitjob, y=tenure)) + geom_boxplot()
```


```
ggplot(df, aes(x=quitjob, y=age)) + geom_boxplot()
```




```
ggplot(df, aes(x=quitjob, y=commute)) + geom_boxplot()
```




```
ggplot(df, aes(x=quitjob, y=costofcommute)) + geom_boxplot()
```




```
# use df with no merged timeserieses to prevent wrong calculations on wages  
# drop not finite numeric values to get rid of warnings  
ggplot(volunteer_endperiod_wage, aes(x=quitjob, y=basewage)) + geom_boxplot()
```



```
ggplot(volunteer_endperiod_wage, aes(x=quitjob, y=bonustotal)) + geom_boxplot()
```




```
ggplot(volunteer_endperiod_wage, aes(x=quitjob, y=grosswage)) + geom_boxplot()
```




```
# Look at correlations
# subselect variable available in both groups (e.g. attitude data not available for quitters)
non_quitter_df <- non_quitter_df %>% select(costofcommute, commute, tenure, age, grosswage, basewage, bonustotal)

quitter_df <- quitter_df %>% select(costofcommute, commute, tenure, age, grosswage, basewage, bonustotal, logd)

nums <- dplyr::select(quitter_df, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
ggcorrplot(cm, lab = FALSE)
```



```
nums <- dplyr::select(non_quitter_df, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
ggcorrplot(cm, lab = FALSE)
```

```
# Make some nice looking box plots for perform, commute
stats <- quitter_df %>% summarise(
  min = min(perform1, na.rm = TRUE),
  max = max(perform1, na.rm = TRUE),
  mean = mean(perform1, na.rm = TRUE),
  median = median(perform1, na.rm = TRUE),
  sd = sd(perform1, na.rm = TRUE),
  n=n()
)
df<-df %>% filter(is.finite(perform1))
p <- ggplot(df, aes(y = perform1, x=quitjob)) +
  geom_boxplot(
    outlier.alpha = 0.6,
    outlier.size = 1.8,
    fill = "#9df3c4",
    color = "#0d0c3f"
  ) +
  labs(y = "Performance score", title = "Performance - quitters vs. non-quitters") +
  theme_minimal()


# add annotations
p + annotate("text", x = 2.3, y = 4,
  label = paste0("SD: ", round(stats$sd,2)),
  hjust = 0, color = "#0d0c3f")+
  annotate("text", x = 2.3, y = 3.6,
  label = paste0("Mean: ", round(stats$mean,2)),
  hjust = 0, color = "#0d0c3f")
```

```

annotate("text", x = 2.3, y = 3.2,
  label = paste0("Median: ", round(stats$median,2)),
  hjust = 0, color = "#0d0c3f") +
annotate("text", x = 2.3, y = 2.8,
  label = paste0("Min: ", round(stats$min,2)),
  hjust = 0, color = "#0d0c3f") +
annotate("text", x = 2.3, y = 2.4,
  label = paste0("Max: ", round(stats$max,2)),
  hjust = 0, color = "#0d0c3f") +
annotate("text", x = 0.27, y = 9000,
  label = paste0("N: ", round(stats$n,2)),
  hjust = 0, color = "#0d0c3f")

```

Performance – quitters vs. non-quitters



```

ggsave("performance_quitter_non_quitters.png", width = 7, height = 5, dpi = 300)

```

```

stats <- quitter_df %>% summarise(
  min = min(commute, na.rm = TRUE),
  max = max(commute, na.rm = TRUE),
  mean = mean(commute, na.rm = TRUE),
  median = median(commute, na.rm = TRUE),
  sd = sd(commute, na.rm = TRUE),
  n=n()
)

p <- ggplot(df, aes(y = commute, x=quitjob)) +
  geom_boxplot(
    outlier.alpha = 0.6,


```

```

    outlier.size = 1.8,
    fill = "#9df3c4",
    color = "#0d0c3f"
  ) +
  labs(y = "Commute score", title = "Commute - quitters vs. non-quitters") +
  theme_minimal()

# add annotations
p + annotate("text", x = 2.3, y = 4,
            label = paste0("SD: ", round(stats$sd,2)),
            hjust = 0, color = "#0d0c3f")+
  annotate("text", x = 2.3, y = 3.6,
            label = paste0("Mean: ", round(stats$mean,2)),
            hjust = 0, color = "#0d0c3f")+
  annotate("text", x = 2.3, y = 3.2,
            label = paste0("Median: ", round(stats$median,2)),
            hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 2.3, y = 2.8,
            label = paste0("Min: ", round(stats$min,2)),
            hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 2.3, y = 2.4,
            label = paste0("Max: ", round(stats$max,2)),
            hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 0.27, y = 9000,
            label = paste0("N: ", round(stats$n,2)),
            hjust = 0, color = "#0d0c3f")

```



```

ggsave("commute_quitter_non_quitters.png", width = 7, height = 5, dpi = 300)


stats <- quitter_df %>% summarise(
  min = min(grosswage, na.rm = TRUE),
  max = max(grosswage, na.rm = TRUE),
  mean = mean(grosswage, na.rm = TRUE),
  median = median(grosswage, na.rm = TRUE),
  sd = sd(grosswage, na.rm = TRUE),
  n=n()
)

p <- ggplot(df, aes(y = grosswage, x=quitjob)) +
  geom_boxplot(
    outlier.alpha = 0.6,
    outlier.size = 1.8,
    fill = "#9df3c4",
    color = "#0d0c3f"
  ) +
  labs(y = "Grosswage in CNY", title = "Grosswage - quitters vs. non-quitters") +
  theme_minimal()

# add annotations
p + annotate("text", x = 2.1, y = 14000,
  label = paste0("SD: ", round(stats$sd,0), "CNY"),
  hjust = 0, color = "#0d0c3f")+
  annotate("text", x = 2.1, y = 13000,
  label = paste0("Mean: ", round(stats$mean,0), "CNY"),
  hjust = 0, color = "#0d0c3f")+
  annotate("text", x = 2.1, y = 12000,
  label = paste0("Median: ", round(stats$median,0), "CNY"),
  hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 2.1, y = 11000,
  label = paste0("Min: ", round(stats$min,0), "CNY"),
  hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 2.1, y = 10000,
  label = paste0("Max: ", round(stats$max,0), "CNY"),
  hjust = 0, color = "#0d0c3f") +
  annotate("text", x = 0.27, y = 9000,
  label = paste0("N: ", round(stats$n,2)),
  hjust = 0, color = "#0d0c3f")

## Warning: Removed 34 rows containing non-finite outside the scale range
## (`stat_boxplot()`).

```



```
ggsave("grosswage_quitter_non_quitters.png", width = 7, height = 5, dpi = 300)
```

```
## Warning: Removed 34 rows containing non-finite outside the scale range
## (`stat_boxplot()`).
```

```
# radar plot to compare quitter vs non quitter
```

```
radar_by_group(df, group_col="quitjob", include_group_indices=c(TRUE, FALSE), title="Quitter vs non-quitter")
```

```
## [1] "commute" "costofcommute" "quitjob"
## [4] "gender_num" "married_num" "bedroom"
## [7] "children_num" "high_educ_num" "promote_switch_num"
## [10] "homethatweek_num" "perform1" "grosswage"
## [13] "exhaustion" "tenure"
```

```
## Warning: Using an external vector in selections was deprecated in tidysselect 1.1.0.
```

```
## i Please use `all_of()` or `any_of()` instead.
```

```
## # Was:
```

```
## data %>% select(final_selection)
```

```
##
```

```
## # Now:
```

```
## data %>% select(all_of(final_selection))
```

```
##
```

```
## See <https://tidysselect.r-lib.org/reference/faq-external-vector.html>.
```

```
## This warning is displayed once every 8 hours.
```

```
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

```
## # A tibble: 2 x 14
```

```
## quitjob commute costofcommute gender married bedroom children high_educ
```

```
##   <lgl>      <dbl>          <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 FALSE      105.          6.67 0.539   0.153   0.961   0.101   0.402
## 2 TRUE       108.          7.65 0.442   0.226   1       0.230   0.417
## # i 6 more variables: promote_switch <dbl>, homethatweek <dbl>, perform1 <dbl>,
## #   grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "quitjob"
## [1] TRUE FALSE
## # A tibble: 2 x 14
##   quitjob commute costofcommute gender married bedroom children high_educ
##   <lgl>      <dbl>          <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 FALSE      105.          6.67 0.539   0.153   0.961   0.101   0.402
## 2 TRUE       108.          7.65 0.442   0.226   1       0.230   0.417
## # i 6 more variables: promote_switch <dbl>, homethatweek <dbl>, perform1 <dbl>,
## #   grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "FALSE" "TRUE"

## pdf
##   2
```

Generally speaking quitters perform less, work for the company shorter than average, have an average age, have slightly higher costs of commute while commuting about the same in minutes per week, earn less and work not from home in comparison to workers who stay. When begin older, quitters have a higher tendency to have children, being married and have higher education. People who were older and left at some point generally performed better than younger quitters. People who stay, tend to have lower cost of commute and lower commute times (maybe because they commit to working for the company and move closer to it).

We can already get a clue from this comparison of the differences between quitters and non quitters about the group constellation and personas.

Groups

Lets have a look on all possible combinations of categorical values and which are represented within the volunteer dataset with a threshold of at least 100 people-work-performances per group. We choose this sample size to increase our confidence. Afterwards we will do the same general analysis we did for the whole population for each group.

Important note: Because all quitters do not have attitude data, positive, negative and exhaustion appear to be “average” on the radar plots, but are actually NULL which is caused by the radar plotting function.

Available groups

```
# make it easier to process
df <- volunteer_endperiod_attitude_performance_wage
# define categorical variables
cat_cols <- c("gender", "married", "high_educ", "bedroom", "children", "promote_switch", "quitjob", "homethatweek")

# factors
df <- df %>% mutate(across(all_of(cat_cols), ~ factor(.x)))

# build grid of categories
lvls <- lapply(df[cat_cols], levels)
grid <- tidyr::expand_grid(!!!lvls)

# small helpers
safe_min <- function(x) {
  x <- x[is.finite(x)]
}
```

```

  if (length(x)) min(x) else NA_real_
}
safe_max <- function(x) {
  x <- x[is.finite(x)]
  if (length(x)) max(x) else NA_real_
}
safe_sd <- function(x) {
  x <- x[is.finite(x)]
  if (length(x) > 1) stats::sd(x) else NA_real_
}
safe_se <- function(x) {
  x <- x[is.finite(x)]
  if (length(x) > 1) stats::sd(x)/sqrt(length(x)) else NA_real_
}

# create summary for each category/group of people working at the company
summary_df <- df %>%
  group_by(across(all_of(cat_cols))) %>%
  summarise(
    # add number per category to filter later
    n = n(),
    across(where(is.numeric),
      list(
        mean = ~mean(.x[is.finite(.x)], na.rm = TRUE),
        sd   = ~safe_sd(.x),
        se   = ~safe_se(.x),
        med  = ~median(.x, na.rm = TRUE),
        min  = ~safe_min(.x),
        max  = ~safe_max(.x)
      ),
      .names = "{.col}_{.fn}"
    ),
    .groups = "drop"
  ) %>%
  mutate(group_index = row_number()) %>% # create an index for each group for better handling
  arrange(desc(n))
# add group_index to the big d
group_characteristics <- summary_df %>% select(group_index, gender, married, bedroom, children, high_edu,
volunteer_endperiod_attitude_performance_wage <- merge(volunteer_endperiod_attitude_performance_wage, group_characteristics, by = "group_index")
# filter by number of rows per group
summary_df <- summary_df %>% filter(n>=100)


```

As a result we receive 23 distinct groups out of $8 \times 8 = 64$ possible options where each group represents a work week of an employee:

```

# First we look at the quitters and try to see which are more likely to quit
ggplot(summary_df, aes(x = factor(group_index), y = quitjob_num_mean)) +
  geom_boxplot() +
  labs(x = "Group index", y = "percentage to quit (0-1)") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



```
df <- volunteer_endperiod_attitude_performance_wage
radar_by_group(df, group_col="group_index", include_group_indices=c(7,16,40,48,60), title="Quitter group")
```

```
## [1] "commute"          "costofcommute"    "group_index"
## [4] "gender_num"       "married_num"      "bedroom"
## [7] "children_num"     "high_educ_num"   "promote_switch_num"
## [10] "quitjob_num"      "homethatweek_num" "perform1"
## [13] "grosswage"        "exhaustion"       "tenure"
## # A tibble: 62 x 15
##   group_index commute costofcommute gender married bedroom children high_educ
##   <int>    <dbl>      <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         1    103.        6.45     1       0       0       0       0
## 2         2    180         18     1       0       0       0       0
## 3         3    120         9      1       0       0       0       0
## 4         4    119.        6.71     1       0       1       0       0
## 5         5    92.0        5.12     1       0       1       0       0
## 6         6    131.        7.58     1       0       1       0       0
## 7         7    118.        4.11     1       0       1       0       0
## 8         8     20         0      1       0       1       0       0
## 9         9    112.        3.88     1       0       1       0       0
## 10        10     28         0      1       0       1       0       0
## # i 52 more rows
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "group_index"
## [1] 7 16 40 48 60
## # A tibble: 5 x 15
```



```
##   group_index commute costofcommute gender married bedroom children high_educ
##           <int>    <dbl>         <dbl> <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1           7    118.           4.11     1       0       1       0       0
## 2          16    52.1           3.57     1       0       1       0       1
## 3          40   120.           8.28     0       0       1       0       0
## 4          48   33.5          11.7     0       0       1       0       1
## 5          60   101.          20.5     0       1       1       1       1
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "7" "16" "40" "48" "60"

## pdf
##    2
```

Here we can clearly identify groups 7,16,40,48 and 60 where all people quit. Within the rest of the groups nobody quit.

Analyze quitter groups

```
df <- volunteer_endperiod_attitude_performance_wage
```

Group 7 Persona observations:

female
not married
number of bedrooms: 1
no children
no high education
no promotion
worked for 7.6 months at the company when left (< avg)
20.5 years old (< avg)

This groups shows a very similar correlation matrix in comparison to the overall sample, no specialities.

Interpretation

Looks like a group of female juniors (wage, performance, tenure, age), living not close by, probably the first or second job after high school(?) or even a job for a gap year. Probably left because of pursuing with school or another job for a pay increase.

Recommendation

This group is among the youngest and has a worse performance than the average. Offering working from home/remote working capabilities can be recommended, as these group could for example continue working during university or traveling. When performance increases a raise would be appropriate but now as first action.

```
# select group data
group7 <- df %>% filter(group_index==7)
stargazer(as.data.frame(group7), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## bedroom        357    1.000    0.000      1         1         1
## children        357    0.000    0.000      0         0         0
## high_educ       357    0.000    0.000      0         0         0
## promote_switch  357    0.000    0.000      0         0         0
```

```
## quitjob      357  1.000  0.000    1    1    1
## homethatweek 357  0.000  0.000    0    0    0
## perform1     345 -0.308  0.854  -3.031  -0.204  2.230
## phonecall     320 -0.220  0.983  -3.054  -0.157  5.828
## phonecallraw  315 411.311 137.944    22    419  1,264
## logphonecall  315  5.941  0.443    3.091    6.038  7.142
## logcallpersec 307 -5.217  0.114   -5.560   -5.198  -4.924
## logcalllength 308 11.142  0.444    8.473   11.222  11.900
## logcall_dayworked 308  9.493  0.369    7.087    9.566  10.049
## logdaysworked 357  1.603  0.325    0.000    1.609  1.946
## homethatweek_num 357  0.000  0.000    0    0    0
## basewage     352 1,450.503 161.296   527.070  1,500.000  1,750.000
## grosswage     352 2,281.241 477.664  1,140.000  2,222.000  4,487.660
## bonustotal    352  830.738 457.575    0.000    795.000  2,937.660
## age          357 20.501  1.315    19    21    22
## tenure       357  7.622  2.661     2     9    10
## commute      357 118.263  97.337    20   120   300
## gender_num   357  1.000  0.000     1     1     1
## married_num  357  0.000  0.000     0     0     0
## high_educ_num 357  0.000  0.000     0     0     0
## children_num  357  0.000  0.000     0     0     0
## costofcommute 357  4.113  4.491    0.000    4.545   10.000
## promote_switch_num 357  0.000  0.000     0     0     0
## quitjob_num  357  1.000  0.000     1     1     1
## group_index   357  7.000  0.000     7     7     7
## -----
```

```
# distinct persons
unique(group7 %>%select(personid))
```

```
##      personid
## 1      38862
## 2      40034
## 3      44256
## 11     40346
## 27     42104
## 109    39478
```

```
group7 <- to_monthly(group7)
```

```
## `summarise()` has grouped output by 'personid'. You can override using the
## `.groups` argument.
```

```
stargazer(as.data.frame(group7), type = "text", median = TRUE, header = TRUE)
```


```
##
## =====
## Statistic      N    Mean    St. Dev.    Min    Median    Max
## -----
## costofcommute   88  4.019    4.508    0.000    0.000   10.000
## bedroom         88  1.000    0.000     1     1     1
## promote_switch   88  0.000    0.000     0     0     0
## promote_switch_num 88  0.000    0.000     0     0     0
## quitjob_num      88  1.000    0.000     1     1     1
## commute         88 118.636   95.979    20    120    300
## tenure          88  7.398    2.806     2     9    10
```

```
## age      88 20.545 1.295 19 21 22
## grosswage 87 2,256.619 482.590 1,140.000 2,221.000 4,487.660
## basewage 87 1,445.713 178.094 527.070 1,500.000 1,750.000
## bonustotal 87 810.906 455.085 0.000 747.000 2,937.660
## logdaysworked 88 1.583 0.273 0.520 1.638 1.946
## logcall_dayworked 74 9.470 0.366 7.087 9.540 9.829
## logcallpersec 74 -5.217 0.102 -5.515 -5.219 -5.016
## logphonecall 77 5.921 0.442 3.091 5.996 6.834
## logcalllength 74 11.113 0.422 8.473 11.169 11.699
## phonecallraw 77 411.533 122.225 22.000 412.800 973.000
## phonecall 78 -0.244 0.902 -3.054 -0.216 3.784
## perform1 84 -0.357 0.726 -3.031 -0.226 0.938
## children_num 88 0.000 0.000 0 0 0
## high_educ_num 88 0.000 0.000 0 0 0
## married_num 88 0.000 0.000 0 0 0
## gender_num 88 1.000 0.000 1 1 1
## homethatweek_num 88 0.000 0.000 0 0 0
## quitjob 88 1.000 0.000 1 1 1
## high_educ 88 0.000 0.000 0 0 0
## children 88 0.000 0.000 0 0 0
## homethatweek 88 0.000 0.000 0 0 0
## -----
```

```
nums <- dplyr::select(group7, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
```

```
## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is
## zero
```

```
ggcorrplot(cm, lab = FALSE)
```



```
# radar plot
radar_by_group(df, group_col="group_index", include_group_indices=c(7), title="Group 7 characteristics")

## [1] "commute"          "costofcommute"    "group_index"
## [4] "gender_num"       "married_num"      "bedroom"
## [7] "children_num"     "high_educ_num"   "promote_switch_num"
## [10] "quitjob_num"      "homethatweek_num" "perform1"
## [13] "grosswage"        "exhaustion"       "tenure"
## # A tibble: 62 x 15
##   group_index commute costofcommute gender married bedroom children high_educ
##   <int>    <dbl>         <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         1    103.         6.45     1       0       0       0       0
## 2         2    180          18       1       0       0       0       0
## 3         3    120           9       1       0       0       0       0
## 4         4    119.        6.71     1       0       1       0       0
## 5         5    92.0        5.12     1       0       1       0       0
## 6         6    131.        7.58     1       0       1       0       0
## 7         7    118.        4.11     1       0       1       0       0
## 8         8     20           0       1       0       1       0       0
## 9         9    112.        3.88     1       0       1       0       0
## 10        10     28           0       1       0       1       0       0
## # i 52 more rows
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "group_index"
## [1] 7
## # A tibble: 1 x 15
```

```
##   group_index commute costofcommute gender married bedroom children high_educ
##           <int>   <dbl>         <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1           7    118.         4.11     1      0      1      0      0
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "7"

## pdf
##    2
```

Group 16 Persona observations:

female
not married
number of bedrooms: 1
no children
high education
no promotion
worked for 23.5 months at the company when left (= avg)
23.7 years old (>= avg)

For this group the strong positive correlation between tenure and age stands out. For the full sample it is slightly negative. Commute and age have a much stronger negative correlation

Interpretation

Looks like a group of female juniors (wage, performance, tenure, age), who already worked for almost 2 years, living somewhat central, probably the first or second job college. Decent performance, not a top performer but average. Probably left because they did not receive a promotion and went for a higher paying job as their salary was still below average.

Recommendation

As this group has an average performance while still being young, and already has a higher education degree (so they probably won't leave for education) our client should prioritize this group and try to give employees in this group a raise and a promotion to keep them within the company. Additionally offer working from home as a benefit, though is not the greatest leverage for this group.

```
# select group data
group16 <- df %>% filter(group_index==16)
stargazer(as.data.frame(group16), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## bedroom        338     1.000     0.000        1         1         1
## children        338     0.000     0.000        0         0         0
## high_educ       338     1.000     0.000        1         1         1
## promote_switch  338     0.000     0.000        0         0         0
## quitjob         338     1.000     0.000        1         1         1
## homethatweek    338     0.000     0.000        0         0         0
## perform1        338    -0.0001    1.079     -3.031     0.049     2.863
## phonecall       320    -0.038     1.045     -3.054     0.048     2.463
## phonecallraw    307   454.397   154.865      51        457      910
## logphonecall    306     6.046     0.419      3.932     6.124     6.813
## logcallpersec   310    -5.159     0.141     -5.495    -5.151    -4.624
## logcalllength   310    11.183     0.562      5.489    11.316    11.991
## logcall_dayworked 310     9.483     0.485      3.697     9.552    10.100
```

```
## logdaysworked      337    1.645    0.395    0.000    1.792    1.946
## homethatweek_num    338    0.000    0.000     0      0      0
## basewage            333 1,599.595  93.253    1,450    1,600    1,750
## grosswage          325 2,966.197 824.625 1,655.000 2,713.000 6,153.380
## bonustotal         333 1,371.143 812.945    5.000    1,189.000 4,403.380
## age                338   23.683    1.309     22     23     26
## tenure             338   23.485    6.243     13     24     30
## commute            227   52.106   41.709      2     60    120
## gender_num         338    1.000    0.000      1      1      1
## married_num        338    0.000    0.000      0      0      0
## high_educ_num      338    1.000    0.000      1      1      1
## children_num       338    0.000    0.000      0      0      0
## costofcommute      338    3.565    2.123      0      3      6
## promote_switch_num 338    0.000    0.000      0      0      0
## quitjob_num        338    1.000    0.000      1      1      1
## group_index        338   16.000    0.000     16     16     16
## -----
```

```
# distinct persons
unique(group16 %>%select(personid))
```

```
##      personid
## 1      36494
## 2      37276
## 5      31150
## 20     26634
## 41     29230
## 85     26934
```

```
group16 <- to_monthly(group16)
```

```
## `summarise()` has grouped output by 'personid'. You can override using the
## `.groups` argument.
```

```
stargazer(as.data.frame(group16), type = "text", median = TRUE, header = TRUE)
```


```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute   83    3.602    2.124      0         3         6
## bedroom         83    1.000    0.000      1         1         1
## promote_switch   83    0.000    0.000      0         0         0
## promote_switch_num 83    0.000    0.000      0         0         0
## quitjob_num     83    1.000    0.000      1         1         1
## commute         56   53.429   43.192      2        60       120
## tenure          83   23.602    6.334     13        24       30
## age             83   23.699    1.323     22        23       26
## grosswage       80 2,950.480 863.340 1,655.000 2,703.015 6,153.380
## basewage        82 1,603.939  94.258    1,450     1,600     1,750
## bonustotal      82 1,351.534 852.448    5.000   1,135.450 4,403.380
## logdaysworked  83    1.598    0.370     0.000     1.645     1.946
## logcall_dayworked 79    9.413    0.711     3.697     9.545    10.002
## logcallpersec   79   -5.155    0.144    -5.489    -5.154    -4.624
## logphonecall    78    6.012    0.347     4.043     6.094     6.561
## logcalllength   79   11.085    0.761     5.489    11.240    11.695
```

```
## phonecallraw      78  444.169  117.239   57.000   453.667  712.000
## phonecall        80  -0.126   0.828   -2.565   -0.015   1.427
## perform1         83  -0.087   0.855   -2.639   -0.081   1.506
## children_num     83   0.000   0.000    0        0        0
## high_educ_num    83   1.000   0.000    1        1        1
## married_num      83   0.000   0.000    0        0        0
## gender_num       83   1.000   0.000    1        1        1
## homethatweek_num 83   0.000   0.000    0        0        0
## quitjob          83   1.000   0.000    1        1        1
## high_educ        83   1.000   0.000    1        1        1
## children         83   0.000   0.000    0        0        0
## homethatweek     83   0.000   0.000    0        0        0
## -----
```

```
nums <- dplyr::select(group16, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
```

```
## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is
## zero
```

```
ggcorrplot(cm, lab = FALSE)
```



```
# radar plot
```

```
radar_by_group(df, group_col="group_index", include_group_indices=c(16), title="Group 16 characteristics")
```

```
## [1] "commute"          "costofcommute"    "group_index"
## [4] "gender_num"       "married_num"      "bedroom"
## [7] "children_num"     "high_educ_num"    "promote_switch_num"
## [10] "quitjob_num"      "homethatweek_num" "perform1"
```

```
## [13] "grosswage"          "exhaustion"          "tenure"
## # A tibble: 62 x 15
##   group_index commute costofcommute gender married bedroom children high_educ
##   <int>    <dbl>         <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         1    103.          6.45     1       0       0       0       0
## 2         2    180          18      1       0       0       0       0
## 3         3    120           9      1       0       0       0       0
## 4         4    119.         6.71     1       0       1       0       0
## 5         5    92.0         5.12     1       0       1       0       0
## 6         6    131.         7.58     1       0       1       0       0
## 7         7    118.         4.11     1       0       1       0       0
## 8         8     20           0      1       0       1       0       0
## 9         9    112.         3.88     1       0       1       0       0
## 10        10     28           0      1       0       1       0       0
## # i 52 more rows
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "group_index"
## [1] 16
## # A tibble: 1 x 15
##   group_index commute costofcommute gender married bedroom children high_educ
##   <int>    <dbl>         <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1        16    52.1         3.57     1       0       1       0       1
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "16"

## pdf
## 2
```

Group 40 Persona observations:

male
not married
number of bedrooms: 1
no children
no high education
no promotion
worked for 26.8 months at the company when left (= avg)
22.8 years old (<= avg)

This correlation matrix is again very similar to the full sample matrix, no outstanding differences.

Interpretation

Young male group, with no high education and family, which has never worked from home. Heavily underperformed, while working as many days as the average. Lives somewhat far away (over avg) with avg cost for commuting. Earns about 400 less than average despite the bad performance. Group7 had similar circumstances, despite commute, but a much better performance with lower salary. But the group is twice as big as this group.

Recommendation

As the performance is worse than average, trying to keep this group of workers should not be prioritized.

```
# select group data
group40 <- df %>% filter(group_index==40)
stargazer(as.data.frame(group40), type = "text", median = TRUE, header = TRUE)
```



```
##
## =====
## Statistic      N      Mean      St. Dev.      Min      Median      Max
## -----
## bedroom        643      1.000      0.000          1          1          1
## children        643      0.000      0.000          0          0          0
## high_educ        643      0.000      0.000          0          0          0
## promote_switch   643      0.000      0.000          0          0          0
## quitjob          643      1.000      0.000          1          1          1
## homethatweek     643      0.000      0.000          0          0          0
## perform1         642     -0.462      0.899     -3.024     -0.446      2.540
## phonecall        638     -0.392      0.900     -3.106     -0.377      2.554
## phonecallraw     631    391.097    135.837          1        385        883
## logphonecall     631      5.880      0.533          0          5.953      6.783
## logcallpersec    629     -5.157      0.174     -5.609     -5.162     -3.178
## logcalllength    630     11.038      0.572          3.178     11.131     11.824
## logcall_dayworked 630      9.327      0.503          2.485      9.406     10.021
## logdaysworked   642      1.697      0.265          0.000      1.792      1.946
## homethatweek_num 643      0.000      0.000          0          0          0
## basewage         633    1,490.934    142.390     650.000    1,500.000    1,900.000
## grosswage        633    2,643.880    737.111    1,360.000    2,485.170    5,011.380
## bonustotal       633    1,152.946    685.716          0.000      966.000    3,461.380
## age              643     22.879      2.293          19         23         27
## tenure           643     26.818      45.869          3         10        150
## commute          612    119.649     51.741          30        120        200
## gender_num       643      0.000      0.000          0          0          0
## married_num      643      0.000      0.000          0          0          0
## high_educ_num    643      0.000      0.000          0          0          0
## children_num     643      0.000      0.000          0          0          0
## costofcommute    643      8.283      2.400          4         10         12
## promote_switch_num 643      0.000      0.000          0          0          0
## quitjob_num      643      1.000      0.000          1          1          1
## group_index      643     40.000      0.000          40         40         40
## -----
```

```
# distinct persons
unique(group40 %>%select(personid))
```

```
##      personid
## 1      44794
## 2      43288
## 4      38552
## 16     38842
## 19     41332
## 55     40008
## 128    43524
## 136    35822
## 177    34890
## 184    41320
## 275    39096
## 406    42096
```

```
group40 <- to_monthly(group40)
```

```
## `summarise()` has grouped output by 'personid'. You can override using the
```

```
## `.groups` argument.
```

```
stargazer(as.data.frame(group40), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute  155    8.284    2.395       4        10        12
## bedroom        155    1.000    0.000       1         1         1
## promote_switch  155    0.000    0.000       0         0         0
## promote_switch_num 155    0.000    0.000       0         0         0
## quitjob_num     155    1.000    0.000       1         1         1
## commute        146   119.041   51.716      30       120       200
## tenure         155    26.045   45.479       3        10       150
## age            155    22.845    2.260      19        23        27
## grosswage       152  2,613.481 734.519  1,360.000 2,440.345 5,011.380
## basewage       152  1,490.737 138.562   650.000  1,500.000 1,900.000
## bonustotal     152  1,122.743 686.273    0.000    950.105  3,461.380
## logdaysworked  155    1.673    0.240       0         1.705    1.946
## logcall_dayworked 153    9.301    0.355       7.774     9.386    9.743
## logcallpersec   153   -5.159    0.152      -5.548    -5.169   -4.687
## logphonecall    153    5.842    0.429       3.367     5.913    6.567
## logcalllength   153   11.002    0.434       8.830    11.106   11.588
## phonecallraw    153   382.913  108.073   29.000   384.000   725.000
## phonecall       155   -0.458    0.745      -2.735    -0.416    1.607
## perform1       155   -0.523    0.748      -2.766    -0.476    1.749
## children_num    155    0.000    0.000       0         0         0
## high_educ_num   155    0.000    0.000       0         0         0
## married_num     155    0.000    0.000       0         0         0
## gender_num      155    0.000    0.000       0         0         0
## homethatweek_num 155    0.000    0.000       0         0         0
## quitjob         155    1.000    0.000       1         1         1
## high_educ       155    0.000    0.000       0         0         0
## children        155    0.000    0.000       0         0         0
## homethatweek    155    0.000    0.000       0         0         0
## -----
```

```
nums <- dplyr::select(group40, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
```

```
## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is
## zero
```

```
ggcorrplot(cm, lab = FALSE)
```



```
##   group_index commute costofcommute gender married bedroom children high_educ
##           <int>   <dbl>         <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1           40    120.         8.28     0       0       1       0       0
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "40"

## pdf
##    2
```

Group 48 Persona observations:

male
not married
number of bedrooms: 1
no children
high education
no promotion
worked for 15.4 months at the company when left (= avg)
21.8 years old (<= avg)

We can see a strong correlation between costofcommute, tenure and age and a rather negative one between costofcommute, commute, tenure, age and grosswage/bonustotal. But small sample size (less confidence).

Interpretation

Young male group, with no education and no family, which has never worked from home. Heavily underperformed, even more than males with high education who quit, while working as many days as the average. Lives closer to work than the average with slightly higher than avg cost for commuting. Earns about 400 less than average despite the bad performance. Unclear why they left, probably also because they wanted more money/a promotion.

Recommendation

Because of the small sample size confidence is low. Tough, as the performance is bad, trying to keep this group of workers should not be prioritized. The company has to ask itself if it is willing to pay more money for this performance,

```
# select group data
group48 <- df %>% filter(group_index==48)
stargazer(as.data.frame(group48), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N    Mean    St. Dev.    Min    Median    Max
## -----
## bedroom        218    1.000    0.000      1      1      1
## children        218    0.000    0.000      0      0      0
## high_educ       218    1.000    0.000      1      1      1
## promote_switch  218    0.000    0.000      0      0      0
## quitjob         218    1.000    0.000      1      1      1
## homethatweek    218    0.000    0.000      0      0      0
## perform1        218   -0.627    0.690     -2.601   -0.614   1.202
## phonecall       216   -0.446    0.750     -2.618   -0.401   2.029
## phonecallraw    218  365.362  104.638      65    373.5    695
## logphonecall    216    5.845    0.379      4.174    5.923    6.544
## logcallpersec   217   -5.210    0.159     -5.721   -5.193   -4.581
## logcalllength   218   11.054    0.336      9.594   11.143   11.468
```

```
## logcall_dayworked 218 9.379 0.298 8.157 9.422 9.880
## logdaysworked 218 1.675 0.208 0.693 1.792 1.946
## homethatweek_num 218 0.000 0.000 0 0 0
## basewage 218 1,565.138 105.495 1,300 1,600 1,850
## grosswage 218 2,652.406 494.384 1,330.000 2,535.000 3,784.760
## bonustotal 218 1,087.269 438.508 30.000 985.340 2,134.760
## age 218 21.844 1.314 20 22 23
## tenure 218 15.486 12.637 4 4 34
## commute 218 33.482 27.988 1 30 75
## gender_num 218 0.000 0.000 0 0 0
## married_num 218 0.000 0.000 0 0 0
## high_educ_num 218 1.000 0.000 1 1 1
## children_num 218 0.000 0.000 0 0 0
## costofcommute 149 11.664 9.548 2 8 25
## promote_switch_num 218 0.000 0.000 0 0 0
## quitjob_num 218 1.000 0.000 1 1 1
## group_index 218 48.000 0.000 48 48 48
## -----
```

```
# distinct persons
unique(group48 %>%select(personid))
```

```
## personid
## 1 24324
## 2 32804
## 6 42618
## 36 42634
```

```
group48 <- to_monthly(group48)
```

```
## `summarise()` has grouped output by 'personid'. You can override using the
## `.groups` argument.
```

```
stargazer(as.data.frame(group48), type = "text", median = TRUE, header = TRUE)
```


```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute  36  11.194    9.627      2         8        25
## bedroom        55   1.000     0.000      1         1         1
## promote_switch  55   0.000     0.000      0         0         0
## promote_switch_num 55   0.000     0.000      0         0         0
## quitjob_num    55   1.000     0.000      1         1         1
## commute        55  31.618    28.033      1        30        75
## tenure         55  14.491    12.576      4         4        34
## age            55  21.745     1.336     20        22        23
## grosswage      55 2,627.605 547.401 1,330.000 2,535.000 3,784.760
## basewage       55 1,560.000 115.229   1,300     1,600     1,850
## bonustotal     55 1,067.605 478.563   30.000    985.340   2,134.760
## logdaysworked  55   1.677     0.125     1.314     1.701     1.884
## logcall_dayworked 55   9.336     0.292     8.157     9.383     9.692
## logcallpersec   55  -5.228     0.160    -5.693    -5.237    -4.817
## logphonecall    55   5.789     0.390     4.174     5.866     6.239
## logcalllength   55  11.014     0.314     9.767    11.120    11.305
## phonecallraw    55  353.232    90.694    65.000    360.500   522.750
```

```
## phonecall      55 -0.545  0.665  -2.618  -0.521  0.754
## perform1      55 -0.707  0.598  -2.601  -0.665  0.410
## children_num  55  0.000  0.000    0      0      0
## high_educ_num 55  1.000  0.000    1      1      1
## married_num   55  0.000  0.000    0      0      0
## gender_num    55  0.000  0.000    0      0      0
## homethatweek_num 55 0.000  0.000    0      0      0
## quitjob       55  1.000  0.000    1      1      1
## high_educ     55  1.000  0.000    1      1      1
## children      55  0.000  0.000    0      0      0
## homethatweek  55  0.000  0.000    0      0      0
## -----
```

```
nums <- dplyr::select(group48, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
```

```
## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is
## zero
```

```
ggcorrplot(cm, lab = FALSE)
```



```
# radar plot
```

```
radar_by_group(df, group_col="group_index", include_group_indices=c(48), title="Group 48 characteristics")
```

```
## [1] "commute"          "costofcommute"    "group_index"
## [4] "gender_num"       "married_num"      "bedroom"
## [7] "children_num"     "high_educ_num"    "promote_switch_num"
## [10] "quitjob_num"      "homethatweek_num" "perform1"
## [13] "grosswage"        "exhaustion"       "tenure"
```

```
## # A tibble: 62 x 15
##   group_index commute costofcommute gender married bedroom children high_educ
##   <int>    <dbl>      <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         1    103.        6.45     1       0       0       0       0
## 2         2    180         18     1       0       0       0       0
## 3         3    120         9      1       0       0       0       0
## 4         4    119.        6.71     1       0       1       0       0
## 5         5     92.0        5.12     1       0       1       0       0
## 6         6    131.        7.58     1       0       1       0       0
## 7         7    118.        4.11     1       0       1       0       0
## 8         8     20         0      1       0       1       0       0
## 9         9    112.        3.88     1       0       1       0       0
## 10        10     28         0      1       0       1       0       0
## # i 52 more rows
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "group_index"
## [1] 48
## # A tibble: 1 x 15
##   group_index commute costofcommute gender married bedroom children high_educ
##   <int>    <dbl>      <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         48    33.5        11.7     0       0       1       0       1
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "48"

## pdf
## 2
```

Group 60 Persona observations:

male
 married
 number of bedrooms: 1
 children
 high education
 no promotion
 worked for 22 months at the company when left (> avg)
 27.8 years old (> avg)

We can see a strong correlation age, tenure and commute. And a moderate between tenure, bonustotal, grosswage, basewage (steady salary raises(?)) But small sample size (less confidence).

Interpretation

Older than avg. male group, with high education and family, which has never worked from home. Under-performance a bit, while working as a little less days as the average. Earn about the same amount less than what they perform less. High cost of commute and average commute minutes per week. Probably left because of this, and maybe the performance also was affect by that. As the salary was steadily improved while working their, this was probably not the main reason.

Recommendation

Because of the small sample size confidence is low. Tough, as the performance is probably worse because of the high commute cost in combination with having a family. A first important action can be the offering of working from home to reduce negative aspects of commuting and increasing performance. If performance actually rises, salary should be increased as well according to that.

```
# select group data
group60 <- df %>% filter(group_index==60)
stargazer(as.data.frame(group60), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## bedroom        202     1.000     0.000        1         1         1
## children        202     1.000     0.000        1         1         1
## high_educ       202     1.000     0.000        1         1         1
## promote_switch  202     0.000     0.000        0         0         0
## quitjob         202     1.000     0.000        1         1         1
## homethatweek    202     0.000     0.000        0         0         0
## perform1        201    -0.216     1.292     -2.996    -0.008     2.164
## phonecall       202    -0.182     1.343     -3.113     0.077     2.729
## phonecallraw    184   435.120   181.523         2        461       823
## logphonecall    184     5.904     0.790         0.693     6.133     6.713
## logcallpersec   184    -5.074     0.196     -5.578    -5.124    -4.533
## logcalllength   183    10.976     0.754         6.271    11.183    11.880
## logcall_dayworked 183     9.416     0.631         5.578     9.599    10.088
## logdaysworked  202     1.532     0.364         0.000     1.609     1.946
## homethatweek_num 202     0.000     0.000         0         0         0
## basewage        195  1,555.128  109.941     1,400     1,550     1,800
## grosswage       195  2,811.191  652.400    1,820.000  2,549.860  4,434.900
## bonustotal      195  1,256.063  592.845     20.000    1,054.000  2,734.900
## age            202    27.847     1.740         25         28         30
## tenure         202    22.064     16.113         8          9         47
## commute        154   101.494    99.220         30         32        240
## gender_num      202     0.000     0.000         0         0         0
## married_num     202     1.000     0.000         1         1         1
## high_educ_num   202     1.000     0.000         1         1         1
## children_num    202     1.000     0.000         1         1         1
## costofcommute   202    20.515    21.521         2          2         55
## promote_switch_num 202     0.000     0.000         0         0         0
## quitjob_num     202     1.000     0.000         1         1         1
## group_index     202    60.000     0.000         60         60         60
## -----
```

```
# distinct persons
unique(group60 %>%select(personid))
```

```
##   personid
## 1    40174
## 2    16334
## 3    39942
## 28   29808
```

```
group60 <- to_monthly(group60)
```

```
## `summarise()` has grouped output by 'personid'. You can override using the
## `.groups` argument.
```

```
stargazer(as.data.frame(group60), type = "text", median = TRUE, header = TRUE)
```

```
##
```




```
## =====
## Statistic      N      Mean      St. Dev.      Min      Median      Max
## -----
## costofcommute  53  19.642   21.524       2         2         55
## bedroom        53   1.000    0.000       1         1         1
## promote_switch  53   0.000    0.000       0         0         0
## promote_switch_num 53   0.000    0.000       0         0         0
## quitjob_num     53   1.000    0.000       1         1         1
## commute         41  97.268   98.467      30        32        240
## tenure          53  22.283   15.823       8         25         47
## age             53  27.698    1.825      25        28         30
## grosswage       50 2,776.677 652.245  1,820.000 2,492.450 4,434.900
## basewage        50 1,564.000 112.957    1,400     1,550     1,800
## bonustotal      50 1,212.677 606.933    20.000    1,021.515 2,734.900
## logdaysworked  53   1.452    0.421     0.000     1.565     1.823
## logcall_dayworked 47   9.298    0.731     6.372     9.529     9.857
## logcallpersec   47  -5.088    0.215    -5.551    -5.123    -4.581
## logphonecall    47   5.757    0.921     2.485     6.073     6.550
## logcalllength   47  10.846    0.896     7.066    11.143    11.553
## phonecallraw    47 412.049  166.167    12.000    451.000    701.000
## phonecall       53  -0.435    1.312    -3.113    -0.179     1.872
## perform1        53  -0.468    1.260    -2.996    -0.188     1.674
## children_num    53   1.000    0.000       1         1         1
## high_educ_num   53   1.000    0.000       1         1         1
## married_num     53   1.000    0.000       1         1         1
## gender_num      53   0.000    0.000       0         0         0
## homethatweek_num 53   0.000    0.000       0         0         0
## quitjob         53   1.000    0.000       1         1         1
## high_educ       53   1.000    0.000       1         1         1
## children        53   1.000    0.000       1         1         1
## homethatweek    53   0.000    0.000       0         0         0
## -----
```

```
nums <- dplyr::select(group60, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")
```

```
## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is
## zero
```

```
ggcorrplot(cm, lab = FALSE)
```



```
# radar plot
radar_by_group(df, group_col="group_index", include_group_indices=c(60), title="Group 60 characteristics")

## [1] "commute"          "costofcommute"    "group_index"
## [4] "gender_num"       "married_num"      "bedroom"
## [7] "children_num"     "high_educ_num"    "promote_switch_num"
## [10] "quitjob_num"      "homethatweek_num" "perform1"
## [13] "grosswage"        "exhaustion"       "tenure"
## # A tibble: 62 x 15
##   group_index commute costofcommute gender married bedroom children high_educ
##   <int>    <dbl>         <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         1    103.         6.45     1       0       0       0       0
## 2         2    180          18       1       0       0       0       0
## 3         3    120           9       1       0       0       0       0
## 4         4    119.        6.71     1       0       1       0       0
## 5         5    92.0        5.12     1       0       1       0       0
## 6         6    131.        7.58     1       0       1       0       0
## 7         7    118.        4.11     1       0       1       0       0
## 8         8     20           0       1       0       1       0       0
## 9         9    112.        3.88     1       0       1       0       0
## 10        10     28           0       1       0       1       0       0
## # i 52 more rows
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "group_index"
## [1] 60
## # A tibble: 1 x 15
```

```
##   group_index commute costofcommute gender married bedroom children high_educ
##           <int>   <dbl>           <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1           60   101.           20.5    0      1      1      1      1
## # i 7 more variables: promote_switch <dbl>, quitjob <dbl>, homethatweek <dbl>,
## #   perform1 <dbl>, grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "60"

## pdf
## 2
```

Compare the quitter groups

```
df <- volunteer_endperiod_attitude_performance_wage
```

```
# plot characteristic radar plot for quitters (per group and summary)
```

```
quitter_group_vec<-c(7,16,40,48,60)
```

```
radar_by_group(df, group_col="group_index", include_group_indices=quitter_group_vec, title="Quitter pro
```

```
## [1] "commute"           "costofcommute"      "negative"
## [4] "positive"          "group_index"        "gender_num"
## [7] "married_num"       "bedroom"            "children_num"
## [10] "high_educ_num"     "promote_switch_num" "quitjob_num"
## [13] "homethatweek_num"  "perform1"           "grosswage"
## [16] "exhaustion"        "tenure"

## # A tibble: 62 x 17
##   group_index commute costofcommute negative positive gender married bedroom
##       <int>   <dbl>           <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         1    103.           6.45    12.4    29.2     1       0       0
## 2         2    180           18      21.8    20.5     1       0       0
## 3         3    120           9       18.8    23       1       0       0
## 4         4    119.           6.71    22.0    22.4     1       0       1
## 5         5    92.0           5.12    14.5    26.6     1       0       1
## 6         6    131.           7.58    18.2    24.8     1       0       1
## 7         7    118.           4.11    NaN     NaN      1       0       1
## 8         8     20           0       NaN     NaN      1       0       1
## 9         9    112.           3.88    13.8    16.2     1       0       1
## 10        10     28           0       NaN     NaN      1       0       1
## # i 52 more rows
## # i 9 more variables: children <dbl>, high_educ <dbl>, promote_switch <dbl>,
## #   quitjob <dbl>, homethatweek <dbl>, perform1 <dbl>, grosswage <dbl>,
## #   exhaustion <dbl>, tenure <dbl>
## [1] "group_index"
## [1] 7 16 40 48 60
## # A tibble: 5 x 17
##   group_index commute costofcommute negative positive gender married bedroom
##       <int>   <dbl>           <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1         7    118.           4.11    NaN     NaN      1       0       1
## 2        16    52.1           3.57    NaN     NaN      1       0       1
## 3        40    120.           8.28    NaN     NaN      0       0       1
## 4        48    33.5           11.7    NaN     NaN      0       0       1
## 5        60    101.           20.5    NaN     NaN      0       1       1
## # i 9 more variables: children <dbl>, high_educ <dbl>, promote_switch <dbl>,
## #   quitjob <dbl>, homethatweek <dbl>, perform1 <dbl>, grosswage <dbl>,
## #   exhaustion <dbl>, tenure <dbl>
## [1] "7" "16" "40" "48" "60"
```

```
## pdf
## 2
```

```
radar_by_group(df, group_col="quitjob", include_group_indices=c(TRUE), title="General quitter profile",
```

```
## [1] "commute"          "costofcommute"      "quitjob"
## [4] "gender_num"        "married_num"        "bedroom"
## [7] "children_num"      "high_educ_num"      "promote_switch_num"
## [10] "homethatweek_num"  "perform1"           "grosswage"
## [13] "exhaustion"        "tenure"
## # A tibble: 2 x 14
##   quitjob commute costofcommute gender married bedroom children high_educ
##   <lgl>    <dbl>          <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 FALSE    103.          6.59  0.513  0.163  0.966  0.116  0.380
## 2 TRUE     105.          9.14  0.450  0.239  1      0.221  0.461
## # i 6 more variables: promote_switch <dbl>, homethatweek <dbl>, perform1 <dbl>,
## #   grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "quitjob"
## [1] TRUE
## # A tibble: 1 x 14
##   quitjob commute costofcommute gender married bedroom children high_educ
##   <lgl>    <dbl>          <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 TRUE     105.          9.14  0.450  0.239  1      0.221  0.461
## # i 6 more variables: promote_switch <dbl>, homethatweek <dbl>, perform1 <dbl>,
## #   grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "TRUE"
```

```
## pdf
## 2
```

```
# plot characteristic radar plot for non quitters (per group and summary)
# get all available groups
all_groups_vec <- unique(df %>% pull(group_index))
non_quitter_group_vec <- setdiff(all_groups_vec, quitter_group_vec)
```

```
radar_by_group(df, group_col="quitjob", include_group_indices=c(FALSE), title="General non-quitter prof
```

```
## [1] "commute"          "costofcommute"      "quitjob"
## [4] "gender_num"        "married_num"        "bedroom"
## [7] "children_num"      "high_educ_num"      "promote_switch_num"
## [10] "homethatweek_num"  "perform1"           "grosswage"
## [13] "exhaustion"        "tenure"
## # A tibble: 2 x 14
##   quitjob commute costofcommute gender married bedroom children high_educ
##   <lgl>    <dbl>          <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 FALSE    103.          6.59  0.513  0.163  0.966  0.116  0.380
## 2 TRUE     105.          9.14  0.450  0.239  1      0.221  0.461
## # i 6 more variables: promote_switch <dbl>, homethatweek <dbl>, perform1 <dbl>,
## #   grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "quitjob"
## [1] FALSE
## # A tibble: 1 x 14
##   quitjob commute costofcommute gender married bedroom children high_educ
##   <lgl>    <dbl>          <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 FALSE    103.          6.59  0.513  0.163  0.966  0.116  0.380
## # i 6 more variables: promote_switch <dbl>, homethatweek <dbl>, perform1 <dbl>,
```

```
## # grosswage <dbl>, exhaustion <dbl>, tenure <dbl>
## [1] "FALSE"

## pdf
## 2

#radar_by_group(df)
```

Top performing 10%

Select

```
df <- volunteer_endperiod_attitude_performance_wage

df <- to_monthly(df)

## `summarise()` has grouped output by 'personid'. You can override using the
## `.groups` argument.

#df <- na.omit(df$perform1)

# define threshold
threshold <- quantile(df$perform1, 0.9, na.rm = TRUE)

# select data
df <- df %>%
  mutate(isTop10 = if_else(perform1 >= threshold, TRUE, FALSE)) %>% filter(!is.na(isTop10))

top10 <- df %>% filter(isTop10 == TRUE)
bottom90 <- df %>% filter(isTop10 == FALSE)
```

Compare

Here we use the same approach as for the quitters:

```
# compare general statistics
stargazer(as.data.frame(top10), type = "text", median = TRUE, header = TRUE)
```

```
##
## =====
## Statistic      N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute  220    6.320     7.435     0.000     4.000     55.000
## positive      53    26.132     5.524    11.000    28.200     39.000
## negative      53    15.038     4.628     8.000    13.500     25.400
## exhaustion    53     7.632     5.855     0.000     6.600    19.250
## bedroom      239     0.954     0.210      0         1         1
## promote_switch 239     0.209     0.408      0         0         1
## promote_switch_num 239     0.209     0.408      0         0         1
## quitjob_num   239     0.155     0.362      0         0         1
## commute       231   104.576    69.229      2         80        300
## tenure        239   34.824    25.884      2         31        150
## age           239   23.728     5.809      1         23         34
## grosswage     236  4,019.654 1,029.368 2,123.000 3,910.500 8,418.000
## basewage      236  1,692.161  194.532   1,400     1,650     2,300
## bonustotal    236  2,327.493  916.875   623.000  2,187.120 6,268.000
```


```
## logdaysworked      239    1.765    0.119    1.337    1.755    1.946
## logcall_dayworked  239    9.773    0.142    9.159    9.791   10.124
## logcallpersec      239   -5.126    0.128   -5.409   -5.126   -4.672
## logphonecall       239    6.411    0.124    5.964    6.394    6.805
## logcalllength      239   11.537    0.146   10.955   11.541   11.854
## phonecallraw       239  620.717   77.990  479.500  607.250  907.000
## phonecall          239    1.181    0.382    0.391    1.114    2.447
## perform1           239    1.354    0.381    0.898    1.262    2.846
## children_num       239    0.243    0.430     0.000     0.000     1.000
## high_educ_num      239    0.469    0.500     0.000     0.000     1.000
## married_num        239    0.331    0.471     0.000     0.000     1.000
## gender_num         239    0.778    0.416     0.000     1.000     1.000
## homethatweek_num   239    0.278    0.446    0.000    0.000    1.000
## quitjob            239    0.155    0.362     0.000     0.000     1.000
## high_educ          239    0.469    0.500     0.000     0.000     1.000
## children           239    0.243    0.430     0.000     0.000     1.000
## homethatweek       239    0.272    0.446     0.000     0.000     1.000
## isTop10            239    1.000    0.000     1.000     1.000     1.000
## -----
```

```
stargazer(as.data.frame(bottom90), type = "text", median = TRUE, header = TRUE)
```


```
##
## =====
## Statistic          N      Mean    St. Dev.    Min      Median      Max
## -----
## costofcommute      2,024    7.345    7.220     0.000     6.000     55.000
## positive           496    24.040    5.826     8.000    24.000     40.000
## negative           496    16.805    6.185     8.000    16.750     40.000
## exhaustion          496     8.676    7.148     0.000     8.000     36.000
## bedroom            2,151     0.977    0.149      0.000      1.000      1.000
## promote_switch     2,151     0.177    0.381      0.000      0.000      1.000
## promote_switch_num 2,151     0.177    0.381      0.000      0.000      1.000
## quitjob_num        2,151     0.255    0.436      0.000      0.000      1.000
## commute            1,931   103.486   67.509      1.000      80.000     300.000
## tenure             2,151    22.703   24.358      2.000     13.000     150.000
## age                2,151    23.464    4.024      1.000     23.000      35.000
## grosswage          2,113  2,972.133  945.290  1,140.000  2,784.000 14,553.000
## basewage           2,123  1,590.618  172.463   527.070  1,550.000  2,450.000
## bonustotal         2,123  1,382.771  878.442    0.000  1,192.690 12,853.000
## logdaysworked     2,150     1.665    0.223     0.000     1.701     1.946
## logcall_dayworked  2,112     9.418    0.429     2.485     9.494    10.129
## logcallpersec      2,112    -5.175    0.166    -5.693    -5.173    -1.099
## logphonecall       2,112     5.936    0.426     0.000     6.021     6.834
## logcalllength      2,112    11.103    0.482     2.485    11.194    11.741
## phonecallraw       2,112   414.398   99.870     1.000   422.625   973.000
## phonecall          2,130    -0.189    0.707    -3.113    -0.104     3.784
## perform1           2,151    -0.213    0.707    -3.031    -0.127     0.897
## children_num       2,151     0.133    0.340      0.000      0.000      1.000
## high_educ_num      2,151     0.402    0.490      0.000      0.000      1.000
## married_num        2,151     0.171    0.376      0.000      0.000      1.000
## gender_num         2,151     0.467    0.499      0.000      0.000      1.000
## homethatweek_num   2,151     0.191    0.386     0.000     0.000    1.000
## quitjob            2,151     0.255    0.436      0.000      0.000      1.000
## high_educ          2,151     0.402    0.490      0.000      0.000      1.000
```

```
## children      2,151  0.133  0.340    0    0    1
## homethatweek  2,151  0.176  0.381    0    0    1
## isTop10       2,151  0.000  0.000    0    0    0
## -----
```


```
# compare distributions
ggplot(df, aes(x=isTop10, y=perform1)) + geom_boxplot()
```



```
ggplot(df, aes(x=isTop10, y=tenure)) + geom_boxplot()
```




```
ggplot(df, aes(x=isTop10, y=age)) + geom_boxplot()
```


```
ggplot(df, aes(x=isTop10, y=commute)) + geom_boxplot()
```

```
## Warning: Removed 228 rows containing non-finite outside the scale range  
## (`stat_boxplot()`).
```




```
ggplot(df, aes(x=isTop10, y=costofcommute)) + geom_boxplot()
```

```
## Warning: Removed 146 rows containing non-finite outside the scale range  
## (`stat_boxplot()`).
```




```
ggplot(df, aes(x=isTop10, y=basewage)) + geom_boxplot()
```

```
## Warning: Removed 31 rows containing non-finite outside the scale range  
## (`stat_boxplot()`).
```




```
ggplot(df, aes(x=isTop10, y=bonustotal)) + geom_boxplot()
```

```
## Warning: Removed 31 rows containing non-finite outside the scale range  
## (`stat_boxplot()`).
```



```
ggplot(df, aes(x=isTop10, y=grosswage)) + geom_boxplot()
```


```
## Warning: Removed 41 rows containing non-finite outside the scale range  
## (`stat_boxplot()`).
```



```
# create correlation matrices  
nums <- dplyr::select(top10, where(is.numeric))  
cm <- cor(nums, use = "pairwise.complete.obs")
```

```
## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is  
## zero
```


```
ggcorrplot(cm, lab = FALSE)
```



```
nums <- dplyr::select(bottom90, where(is.numeric))
cm <- cor(nums, use = "pairwise.complete.obs")

## Warning in cor(nums, use = "pairwise.complete.obs"): the standard deviation is
## zero

ggcorrplot(cm, lab = FALSE)
```



```
# radar plot
radar_by_group(df, group_col="isTop10", include_group_indices=c(TRUE, FALSE), title="Top vs low performance")
```

```
## [1] "commute" "costofcommute" "negative"
## [4] "positive" "isTop10" "gender_num"
## [7] "married_num" "bedroom" "children_num"
## [10] "high_educ_num" "promote_switch_num" "quitjob_num"
## [13] "homethatweek_num" "perform1" "grosswage"
## [16] "exhaustion" "tenure"
## # A tibble: 2 x 17
## isTop10 commute costofcommute negative positive gender married bedroom
## <lgl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 FALSE 103. 7.35 16.8 24.0 0.467 0.171 0.977
## 2 TRUE 105. 6.32 15.0 26.1 0.778 0.331 0.954
## # i 9 more variables: children <dbl>, high_educ <dbl>, promote_switch <dbl>,
## # quitjob <dbl>, homethatweek <dbl>, perform1 <dbl>, grosswage <dbl>,
## # exhaustion <dbl>, tenure <dbl>
## [1] "isTop10"
## [1] TRUE FALSE
## # A tibble: 2 x 17
## isTop10 commute costofcommute negative positive gender married bedroom
## <lgl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 FALSE 103. 7.35 16.8 24.0 0.467 0.171 0.977
## 2 TRUE 105. 6.32 15.0 26.1 0.778 0.331 0.954
## # i 9 more variables: children <dbl>, high_educ <dbl>, promote_switch <dbl>,
## # quitjob <dbl>, homethatweek <dbl>, perform1 <dbl>, grosswage <dbl>,
## # exhaustion <dbl>, tenure <dbl>
```



```
## [1] "FALSE" "TRUE"
```

```
## pdf
```

```
## 2
```

The average employee among the top 10% performing ones, has no children, though it is more likely than for the other 90%, works 25% of the time from home, is almost 24 years old, works for 32 months at the company and an average commute and average cost of commute. Almost 22% receive a promotion, whereas 17.7% of lower performing people receive one. They earn over the average by bonuses, but have an average base salary. Though some not high performing individuals receive enormous bonuses. Only 18% high performing people quit while 25% quit from the other 90%.

Top performer are less likely to quit when receiving a higher bonus than non performer, but more likely if they have a long commute as this is negatively correlative with their negative attitude. Working from home does not affect performance of both groups. Working from home increases positivity while lowering negativity and exhaustion for the top performer.

Recommendations

- Bonuses combined with goals/milestones to motivate people and increase performance, but only if you reach them
- check existing bonuses for not high performing people and reevaluate if that is fair, this might be a motivation killer
- offer more working from home for high performers, as it does not affect their performance and rather increases positivity

Make perform1 and grosswage poxplots look nice

```
p <- ggplot(df, aes(y = perform1, x=isTop10)) +  
  geom_boxplot(  
    outlier.alpha = 0.6,  
    outlier.size = 1.8,  
    fill = "#9df3c4",  
    color = "#0d0c3f"  
  ) +  
  labs(y = "Performance score", title = "Performance - Top-Performer vs. others") +  
  theme_minimal()  
  
# add annotations  
p
```




```
ggsave("performance_top_low.png", width = 7, height = 5, dpi = 300)
```

```
p <- ggplot(df, aes(y = grosswage, x=isTop10)) +
  geom_boxplot(
    outlier.alpha = 0.6,
    outlier.size = 1.8,
    fill = "#9df3c4",
    color = "#0d0c3f"
  ) +
  labs(y = "Grosswage in CNY", title = "Grosswage - Top-Performer vs. others") +
  theme_minimal()
```

```
# add annotations
```

```
p
```

```
## Warning: Removed 41 rows containing non-finite outside the scale range
## (`stat_boxplot()`).
```



```
colnames(df)
```

```
## [1] "personid"      "year_month"    "costofcommute"
## [4] "positive"      "negative"      "exhaustion"
## [7] "bedroom"      "promote_switch" "promote_switch_num"
## [10] "quitjob_num"   "commute"       "tenure"
## [13] "age"          "grosswage"     "basewage"
## [16] "bonustotal"    "logdaysworked" "logcall_dayworked"
## [19] "logcallpersec" "logphonecall"  "logcalllength"
## [22] "phonecallraw"  "phonecall"     "perform1"
## [25] "children_num"  "high_educ_num" "married_num"
## [28] "gender_num"    "homethatweek_num" "quitjob"
## [31] "married"       "high_educ"     "children"
## [34] "homethatweek"  "gender"        "isTop10"
```

```
ggsave("grosswage_top_low.png", width = 7, height = 5, dpi = 300)
```

```
## Warning: Removed 41 rows containing non-finite outside the scale range
## (`stat_boxplot()`).
```

Hypothesis testing

Quitters vs non quitters

```
# H0: There is no difference in earnings between quitters and non-quitters
# H1: There is a significant difference in earnings between quitters and non-quitters
volunteer_endperiod_wage %>%
```

```

group_by(personid, quitjob_num) %>%
summarise(mean_grossw = mean(grosswage, na.rm = TRUE), .groups = "drop") %>%
filter(!is.na(mean_grossw)) %>%
t.test(mean_grossw ~ quitjob_num, data = ., alternative = "greater")

##
## Welch Two Sample t-test
##
## data: mean_grossw by quitjob_num
## t = 6.4985, df = 95.835, p-value = 1.815e-09
## alternative hypothesis: true difference in means between group 0 and group 1 is greater than 0
## 95 percent confidence interval:
## 509.586 Inf
## sample estimates:
## mean in group 0 mean in group 1
## 3293.432 2608.889
# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no difference between perform1 of quitters and non-quitters
#H1: There is a significant difference between perform1 of quitters and non-quitters
volunteer_endperiod_attitude_performance_wage %>%
group_by(personid, quitjob_num) %>%
summarise(mean_perform1 = mean(perform1, na.rm = TRUE), .groups = "drop") %>%
filter(!is.na(mean_perform1)) %>%
t.test(mean_perform1 ~ quitjob_num, data = ., alternative = "less")

##
## Welch Two Sample t-test
##
## data: mean_perform1 by quitjob_num
## t = 3.0921, df = 76.915, p-value = 0.9986
## alternative hypothesis: true difference in means between group 0 and group 1 is less than 0
## 95 percent confidence interval:
## -Inf 0.4601867
## sample estimates:
## mean in group 0 mean in group 1
## 0.02464293 -0.27448432
# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no difference in tenure between quitters and non-quitters
#H1: There is a significant difference in tenure between quitters and non-quitters
volunteer_endperiod_attitude_performance_wage %>%
group_by(personid, quitjob_num) %>%
summarise(mean_tenure = mean(tenure, na.rm = TRUE), .groups = "drop") %>%
filter(!is.na(mean_tenure)) %>%
t.test(mean_tenure ~ quitjob_num, data = ., alternative = "greater")

##
## Welch Two Sample t-test
##
## data: mean_tenure by quitjob_num

```

```

## t = 0.37039, df = 64.932, p-value = 0.3561
## alternative hypothesis: true difference in means between group 0 and group 1 is greater than 0
## 95 percent confidence interval:
## -6.401416      Inf
## sample estimates:
## mean in group 0 mean in group 1
##      23.92632      22.10000

# we fail to reject the null hypothesis, as p-value is greater than 0.05

#H0: there is no difference in commutecost between quitters and non-quitters
#H1: there is a difference in commutecost between quitters and non-quitters
volunteer_endperiod_attitude_performance_wage %>%
  group_by(personid, quitjob_num) %>%
  summarise(mean_costofcommute = mean(costofcommute, na.rm = TRUE), .groups = "drop") %>%
  filter(!is.na(mean_costofcommute)) %>%
  t.test(mean_costofcommute ~ quitjob_num, data = .)

##
## Welch Two Sample t-test
##
## data: mean_costofcommute by quitjob_num
## t = -1.3774, df = 50.59, p-value = 0.1744
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## -5.597109  1.042489
## sample estimates:
## mean in group 0 mean in group 1
##      6.608471      8.885781

# we fail to reject the null hypothesis, as p-value is greater than 0.05

#H0: there is no difference in commute between quitters and non-quitters
#H1: there is a difference in commute between quitters and non-quitters
volunteer_endperiod_attitude_performance_wage %>%
  group_by(personid, quitjob_num) %>%
  summarise(mean_commute = mean(commute, na.rm = TRUE), .groups = "drop") %>%
  filter(!is.na(mean_commute)) %>%
  t.test(mean_commute ~ quitjob_num, data = .)

##
## Welch Two Sample t-test
##
## data: mean_commute by quitjob_num
## t = -0.097326, df = 55.134, p-value = 0.9228
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## -30.92788  28.06286
## sample estimates:
## mean in group 0 mean in group 1
##      104.3103      105.7429

# we fail to reject the null hypothesis, as p-value is greater than 0.05

```

```

#H0: there is no difference in homeoffice_days between quitters and non-quitters
#H1: there is a significant difference in homeoffice_days between quitters and non-quitters
volunteer_endperiod_attitude_performance_wage %>%
  group_by(personid, quitjob_num) %>%
  summarise(mean_homethatweek = mean(homethatweek_num, na.rm = TRUE), .groups = "drop") %>%
  filter(!is.na(mean_homethatweek)) %>%
  t.test(mean_homethatweek ~ quitjob_num, data = ., alternative = "greater")

##
## Welch Two Sample t-test
##
## data: mean_homethatweek by quitjob_num
## t = 4.9726, df = 113.28, p-value = 1.185e-06
## alternative hypothesis: true difference in means between group 0 and group 1 is greater than 0
## 95 percent confidence interval:
##  0.1012484      Inf
## sample estimates:
## mean in group 0 mean in group 1
##    0.22867153    0.07675795

# we reject the null hypothesis, as p-value is below 0.05

#H0: there is no correlation between age and children_num among quitters
#H1: there is a correlation between age and children_num among quitters
volunteer_endperiod_attitude_performance_wage %>%
  filter(quitjob_num == 1) %>%
  t.test(children_num ~ (age > median(age, na.rm = TRUE)),
    data = ., alternative = "less")

##
## Welch Two Sample t-test
##
## data: children_num by age > median(age, na.rm = TRUE)
## t = -27.221, df = 1135.4, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -0.4341011
## sample estimates:
## mean in group FALSE mean in group TRUE
##    0.03795721    0.50000000

# we reject the null hypothesis, as p-value is below 0.05

#H0: there is no correlation between age and married_num among quitters
#H1: there is a correlation between age and married_num among quitters
volunteer_endperiod_attitude_performance_wage %>%
  filter(quitjob_num == 1) %>%
  t.test(married_num ~ (age > median(age, na.rm = TRUE)),
    data = ., alternative = "less")

##
## Welch Two Sample t-test
##

```

```

## data: married_num by age > median(age, na.rm = TRUE)
## t = -17.35, df = 1467.5, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -0.2842208
## sample estimates:
## mean in group FALSE mean in group TRUE
##      0.1145618      0.4285714
# we reject the null hypothesis, as p-value is below 0.05

#H0: there is no correlation between age and high_educ_num among quitters
#H1: there is a correlation between age and high_educ_num among quitters
volunteer_endperiod_attitude_performance_wage %>%
  filter(quitjob_num == 1) %>%
  t.test(high_educ_num ~ (age > median(age, na.rm = TRUE)),
    data = ., alternative = "less")

##
## Welch Two Sample t-test
##
## data: high_educ_num by age > median(age, na.rm = TRUE)
## t = -15.705, df = 2020.9, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -0.2790255
## sample estimates:
## mean in group FALSE mean in group TRUE
##      0.3374741      0.6491597
# we reject the null hypothesis, as p-value is below 0.05

#H0: there is no correlation between mean(perform1) and age among quitters
#H1: there is a correlation between mean(perform1) and age among quitters
volunteer_endperiod_attitude_performance_wage %>%
  filter(quitjob_num == 1) %>%
  t.test(perform1 ~ (age > median(age, na.rm = TRUE)),
    data = ., alternative = "less")

##
## Welch Two Sample t-test
##
## data: perform1 by age > median(age, na.rm = TRUE)
## t = -6.0454, df = 1700.1, p-value = 9.138e-10
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -0.1953103
## sample estimates:
## mean in group FALSE mean in group TRUE
##      -0.34598213      -0.07761384
# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no relationship between commute and quitjob_num

```

```
#H1: Higher commute leads to changes quitjob_num
df <- volunteer_endperiod_attitude_performance_wage
t.test(df$quitjob_num ~ (df$commute > median(df$commute, na.rm = TRUE)),
       data = df)
```

```
##
## Welch Two Sample t-test
##
## data: df$quitjob_num by df$commute > median(df$commute, na.rm = TRUE)
## t = -7.6378, df = 8905.3, p-value = 2.437e-14
## alternative hypothesis: true difference in means between group FALSE and group TRUE is not equal to 0
## 95 percent confidence interval:
## -0.08488709 -0.05021378
## sample estimates:
## mean in group FALSE mean in group TRUE
## 0.1994016 0.2669520
```

```
# we fail to reject the null hypothesis, as p-value is above 0.05
```

Group Hypotheses

Group 7

```
#H0: There is no difference between group 7 mean(quitjob_num) compared to overall average
#H1: Group 7 has higher mean(quitjob_num) compared to overall average
```

```
t.test(group7$quitjob_num, volunteer_endperiod_attitude_performance_wage$quitjob_num, alternative = "greater")
```

```
##
## Welch Two Sample t-test
##
## data: group7$quitjob_num and volunteer_endperiod_attitude_performance_wage$quitjob_num
## t = 179.52, df = 10078, p-value < 2.2e-16
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## 0.7548015 Inf
## sample estimates:
## mean of x mean of y
## 1.0000000 0.2382181
```

```
# we reject the null hypothesis, as p-value is below 0.05
```

```
#H0: There is no difference between group 7 mean(perform1) compared to overall average
#H1: Group 7 has lower mean(perform1) compared to overall average
```

```
t.test(group7$perform1, volunteer_endperiod_attitude_performance_wage$perform1, alternative = "less")
```

```
##
## Welch Two Sample t-test
##
## data: group7$perform1 and volunteer_endperiod_attitude_performance_wage$perform1
## t = -4.2834, df = 85.646, p-value = 2.395e-05
## alternative hypothesis: true difference in means is less than 0
## 95 percent confidence interval:
## -Inf -0.2092549
## sample estimates:
## mean of x mean of y
```



```
## -0.35702606 -0.01499374
```

```
# we reject the null hypothesis, as p-value is below 0.05
```

Group 16

```
#H0: There is no difference between group 16 mean(quitjob_num) compared to overall average
```

```
#H1: Group 16 has higher mean(quitjob_num) compared to overall average
```

```
t.test(group16$quitjob_num, volunteer_endperiod_attitude_performance_wage$quitjob_num, alternative = "g
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: group16$quitjob_num and volunteer_endperiod_attitude_performance_wage$quitjob_num
```

```
## t = 179.52, df = 10078, p-value < 2.2e-16
```

```
## alternative hypothesis: true difference in means is greater than 0
```

```
## 95 percent confidence interval:
```

```
## 0.7548015 Inf
```

```
## sample estimates:
```

```
## mean of x mean of y
```

```
## 1.0000000 0.2382181
```

```
# we reject the null hypothesis, as p-value is below 0.05
```

```
#H0: There is no difference between group 16 mean(perform1) compared to overall average
```

```
#H1: Group 16 has lower mean(perform1) compared to overall average
```

```
t.test(group16$perform1, volunteer_endperiod_attitude_performance_wage$perform1, alternative = "less")
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: group16$perform1 and volunteer_endperiod_attitude_performance_wage$perform1
```

```
## t = -0.7665, df = 83.857, p-value = 0.2228
```

```
## alternative hypothesis: true difference in means is less than 0
```

```
## 95 percent confidence interval:
```

```
## -Inf 0.08467575
```

```
## sample estimates:
```

```
## mean of x mean of y
```

```
## -0.08737226 -0.01499374
```

```
# we fail to reject the null hypothesis, as p-value is above 0.05
```

```
#H0: There is no correlation between age and tenure
```

```
#H1: There is a correlation between age and tenure
```

```
t.test(group16$tenure ~ (group16$age > median(group16$age, na.rm = TRUE)),  
data = group16, alternative = "less")
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: group16$tenure by group16$age > median(group16$age, na.rm = TRUE)
```

```
## t = -13.024, df = 53.123, p-value < 2.2e-16
```

```
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
```

```
## 95 percent confidence interval:
```

```
## -Inf -9.001616
```

```
## sample estimates:
```

```
## mean in group FALSE mean in group TRUE
##          18.50000          28.82927
# we reject the null hypothesis, as p-value is below 0.05
```

Group 40

```
#H0: There is no difference between group 40 mean(quitjob_num) compared to overall average
#H1: Group 40 has higher mean(quitjob_num) compared to overall average
t.test(group40$quitjob_num, volunteer_endperiod_attitude_performance_wage$quitjob_num, alternative = "g")

##
## Welch Two Sample t-test
##
## data: group40$quitjob_num and volunteer_endperiod_attitude_performance_wage$quitjob_num
## t = 179.52, df = 10078, p-value < 2.2e-16
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
##  0.7548015      Inf
## sample estimates:
## mean of x mean of y
## 1.0000000 0.2382181
# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no difference between group 40 mean(perform1) compared to overall average
#H1: Group 40 has lower mean(perform1) compared to overall average
t.test(group40$perform1, volunteer_endperiod_attitude_performance_wage$perform1, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group40$perform1 and volunteer_endperiod_attitude_performance_wage$perform1
## t = -8.3423, df = 162.59, p-value = 1.483e-14
## alternative hypothesis: true difference in means is less than 0
## 95 percent confidence interval:
##      -Inf -0.4072643
## sample estimates:
## mean of x mean of y
## -0.52299524 -0.01499374
# we reject the null hypothesis, as p-value is below 0.05
```

Group 48

```
#H0: There is no difference between group 48 mean(quitjob_num) compared to overall average
#H1: Group 48 has higher mean(quitjob_num) compared to overall average
t.test(group48$quitjob_num, volunteer_endperiod_attitude_performance_wage$quitjob_num, alternative = "g")

##
## Welch Two Sample t-test
##
## data: group48$quitjob_num and volunteer_endperiod_attitude_performance_wage$quitjob_num
## t = 179.52, df = 10078, p-value < 2.2e-16
## alternative hypothesis: true difference in means is greater than 0
```

```
## 95 percent confidence interval:
## 0.7548015      Inf
## sample estimates:
## mean of x mean of y
## 1.0000000 0.2382181

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no difference between group 48 mean(perform1) compared to overall average
#H1: Group 48 has lower mean(perform1) compared to overall average
t.test(group48$perform1, volunteer_endperiod_attitude_performance_wage$perform1, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group48$perform1 and volunteer_endperiod_attitude_performance_wage$perform1
## t = -8.518, df = 55.661, p-value = 5.735e-12
## alternative hypothesis: true difference in means is less than 0
## 95 percent confidence interval:
##      -Inf -0.5561751
## sample estimates:
## mean of x mean of y
## -0.70707354 -0.01499374

# we reject the null hypothesis, as p-value is below 0.05
```

Group 60

```
#H0: There is no difference between group 60 mean(quitjob_num) compared to overall average
#H1: Group 60 has higher mean(quitjob_num) compared to overall average
t.test(group60$quitjob_num, volunteer_endperiod_attitude_performance_wage$quitjob_num, alternative = "greater")

##
## Welch Two Sample t-test
##
## data: group60$quitjob_num and volunteer_endperiod_attitude_performance_wage$quitjob_num
## t = 179.52, df = 10078, p-value < 2.2e-16
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## 0.7548015      Inf
## sample estimates:
## mean of x mean of y
## 1.0000000 0.2382181

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no difference between group 60 mean(perform1) compared to overall average
#H1: Group 60 has lower mean(perform1) compared to overall average
t.test(group60$perform1, volunteer_endperiod_attitude_performance_wage$perform1, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group60$perform1 and volunteer_endperiod_attitude_performance_wage$perform1
## t = -2.6122, df = 52.345, p-value = 0.005857
## alternative hypothesis: true difference in means is less than 0
```

```

## 95 percent confidence interval:
##      -Inf -0.1625289
## sample estimates:
## mean of x mean of y
## -0.46776455 -0.01499374

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no relation between age and commute
#H1: Older Members have higher commute
t.test(group60$commute ~ (group60$age > median(group60$age, na.rm = TRUE)),
       data = group60, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group60$commute by group60$age > median(group60$age, na.rm = TRUE)
## t = -1086, df = 27, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -208.6722
## sample estimates:
## mean in group FALSE mean in group TRUE
##           31           240

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no relation between age and tenure
#H1: Older Members have higher tenure
t.test(group60$tenure ~ (group60$age > median(group60$age, na.rm = TRUE)),
       data = group60, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group60$tenure by group60$age > median(group60$age, na.rm = TRUE)
## t = -25.894, df = 39, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -30.619
## sample estimates:
## mean in group FALSE mean in group TRUE
##           14.25           47.00

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no relation between tenure on grosswage
#H1: Higher tenure leads to higher grosswage
t.test(group60$grosswage ~ (group60$tenure > median(group60$tenure, na.rm = TRUE)),
       data = group60, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group60$grosswage by group60$tenure > median(group60$tenure, na.rm = TRUE)
## t = -6.0257, df = 40.168, p-value = 2.141e-07

```

```

## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -570.2138
## sample estimates:
## mean in group FALSE  mean in group TRUE
##      2586.760      3378.081

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no relation between tenure on basewage
#H1: Higher tenure leads to higher basewage
t.test(group60$basewage ~ (group60$tenure > median(group60$tenure, na.rm = TRUE)),
       data = group60, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group60$basewage by group60$tenure > median(group60$tenure, na.rm = TRUE)
## t = -4.7499, df = 41.164, p-value = 1.24e-05
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -73.07008
## sample estimates:
## mean in group FALSE  mean in group TRUE
##      1536.842      1650.000

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no relation between tenure on bonustotal
#H1: Higher tenure leads to higher bonustotal
t.test(group60$bonustotal ~ (group60$tenure > median(group60$tenure, na.rm = TRUE)),
       data = group60, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group60$bonustotal by group60$tenure > median(group60$tenure, na.rm = TRUE)
## t = -5.1879, df = 36.347, p-value = 4.124e-06
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -457.5232
## sample estimates:
## mean in group FALSE  mean in group TRUE
##      1049.918      1728.081

# we reject the null hypothesis, as p-value is below 0.05

#H0: There is no relationship between commute and perform1
#H1: Higher commute leads to lower perform1
t.test(group60$perform1 ~ (group60$commute > median(group60$commute, na.rm = TRUE)),
       data = group60, alternative = "less")

##
## Welch Two Sample t-test
##
## data: group60$perform1 by group60$commute > median(group60$commute, na.rm = TRUE)

```

```
## t = -3.4552, df = 23.723, p-value = 0.001041
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -0.6877977
## sample estimates:
## mean in group FALSE  mean in group TRUE
##      -0.7907757      0.5722787
# we reject the null hypothesis, as p-value is below 0.05
```

Top 10 performers

```
#H0: The variance(basewage) among top performers is equal to the variance(basewage) among the rest
#H1: The variance(basewage) among top performers is ~ 0
var.test(top10$basewage, bottom90$basewage)
```

```
##
## F test to compare two variances
##
## data: top10$basewage and bottom90$basewage
## F = 1.2723, num df = 235, denom df = 2122, p-value = 0.009799
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
##  1.058939 1.551378
## sample estimates:
## ratio of variances
##      1.272293
```

```
# we reject the null hypothesis, as p-value is below 0.05
```

```
#H0: The mean(grosswage) of Top 10 is equal to the mean(grosswage) of the rest
#H1: The mean(grosswage) of Top 10 is greater than the mean(grosswage) of the rest
t.test(top10$grosswage, bottom90$grosswage, alternative = "greater")
```

```
##
## Welch Two Sample t-test
##
## data: top10$grosswage and bottom90$grosswage
## t = 14.945, df = 281.08, p-value < 2.2e-16
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
##  931.8502      Inf
## sample estimates:
## mean of x mean of y
##  4019.654  2972.133
```

```
# we reject the null hypothesis, as p-value is below 0.05
```

```
#H0 the mean(children_num) of Top 10 is equal to the mean(children_num) of the rest
#H1: The mean(children_num) of Top 10 is higher than the mean(children_num) of the rest
t.test(top10$children_num, bottom90$children_num, alternative = "greater")
```

```
##
## Welch Two Sample t-test
##
## data: top10$children_num and bottom90$children_num
```

```

## t = 3.8014, df = 272.16, p-value = 8.881e-05
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## 0.06181701      Inf
## sample estimates:
## mean of x mean of y
## 0.2426778 0.1334263

# we reject the null hypothesis, as p-value is below 0.05

#H0: The mean(homethatweek_num) of Top 10 is equal to the mean(homethatweek_num) of the rest
#H1: The mean(homethatweek_num) of Top 10 is = 0.25
t.test(top10$homethatweek_num, bottom90$homethatweek_num, alternative = "greater")

##
## Welch Two Sample t-test
##
## data: top10$homethatweek_num and bottom90$homethatweek_num
## t = 2.9062, df = 278.9, p-value = 0.001976
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## 0.03773537      Inf
## sample estimates:
## mean of x mean of y
## 0.2781729 0.1908492

# we reject the null hypothesis, as p-value is below 0.05
# result for top 10 is ~ 0.25 (25.4)

#H0: The mean(age) of top performs is equal to the mean(age) of the rest
#H1: The mean(age) of top performs = 24
t.test(top10$age, bottom90$age, alternative = "greater")

##
## Welch Two Sample t-test
##
## data: top10$age and bottom90$age
## t = 0.68353, df = 263.97, p-value = 0.2474
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## -0.3729608      Inf
## sample estimates:
## mean of x mean of y
## 23.72803 23.46444

#H0: The mean(positivity) of top performs is equal to the mean(positivity) of the rest
#H1: The mean(positivity) of top performs greater than the mean(positivity) of the rest
t.test(top10$positive, bottom90$positive, alternative = "greater")

##
## Welch Two Sample t-test
##
## data: top10$positive and bottom90$positive
## t = 2.6066, df = 64.998, p-value = 0.005664
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:

```

```
## 0.7528422      Inf
## sample estimates:
## mean of x mean of y
## 26.13236 24.04019
```

```
# we reject the null hypothesis, as p-value is below 0.05
# result for top 10 is ~ 24 (23.8)
```

```
#H0: The mean(tenure) of top performs is equal to the mean(tenure) of the rest
#H1: The mean(tenure) of top performs > mean(tenure) of the rest
t.test(top10$tenure, bottom90$tenure, alternative = "greater")
```

```
##
## Welch Two Sample t-test
##
## data: top10$tenure and bottom90$tenure
## t = 6.9078, df = 286.84, p-value = 1.589e-11
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## 9.225705      Inf
## sample estimates:
## mean of x mean of y
## 34.82427 22.70293
```

```
# we reject the null hypothesis, as p-value is below 0.05
```

```
#H0: the mean(promote_switch_num) of top performs is equal to the mean(promote_switch_num) of the rest
#H1: The mean(promote_switch_num) of top performs > mean(promote_switch_num) of the rest
t.test(top10$promote_switch_num, bottom90$promote_switch_num, alternative = "greater")
```

```
##
## Welch Two Sample t-test
##
## data: top10$promote_switch_num and bottom90$promote_switch_num
## t = 1.1783, df = 286.28, p-value = 0.1198
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
## -0.01303256      Inf
## sample estimates:
## mean of x mean of y
## 0.209205 0.176662
```

```
# we reject the null hypothesis, as p-value is below 0.05
```

```
#H0: The mean (quitjob_num) of top performs is equal to the mean(quitjob_num) of the rest
#H1: The mean(quitjob_num) of top performs < mean(quitjob_num) of the rest
t.test(top10$quitjob_num, bottom90$quitjob_num, alternative = "less")
```

```
##
## Welch Two Sample t-test
##
## data: top10$quitjob_num and bottom90$quitjob_num
## t = -3.975, df = 319.79, p-value = 4.352e-05
## alternative hypothesis: true difference in means is less than 0
## 95 percent confidence interval:
##      -Inf -0.05874496
```



```

## sample estimates:
## mean of x mean of y
## 0.1548117 0.2552301

# we reject the null hypothesis, as p-value is below 0.05

#H0: Among top vs. rest, mean(bonustotal) has no effect on quit probability
#H1: Top performers with higher mean(bonustotal) are less likely to quit than non-performers
t.test(top10$bonustotal ~ (top10$bonustotal > median(top10$bonustotal, na.rm = TRUE)),
       data = top10, alternative = "less")

##
## Welch Two Sample t-test
##
## data: top10$bonustotal by top10$bonustotal > median(top10$bonustotal, na.rm = TRUE)
## t = -20.054, df = 179.52, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group FALSE and group TRUE is less than 0
## 95 percent confidence interval:
##      -Inf -1334.946
## sample estimates:
## mean in group FALSE mean in group TRUE
##      1600.043      3054.942

# we reject the null hypothesis, as p-value is below 0.05

#H0: mean(positive) does not differ from homethatweek_num
#H1: higher homethatweek_num leads to higher mean(positive)
t.test(top10$positive ~ (top10$homethatweek_num > median(top10$homethatweek_num, na.rm = TRUE)),
       data = top10, alternative = "greater")

##
## Welch Two Sample t-test
##
## data: top10$positive by top10$homethatweek_num > median(top10$homethatweek_num, na.rm = TRUE)
## t = 3.4124, df = 46.632, p-value = 0.0006701
## alternative hypothesis: true difference in means between group FALSE and group TRUE is greater than 0
## 95 percent confidence interval:
##  2.027312      Inf
## sample estimates:
## mean in group FALSE mean in group TRUE
##  28.76667      24.77757

# we reject the null hypothesis, as p-value is below 0.05

#H0: mean(negative) does not differ from homethatweek_num
#H1: higher homethatweek_num leads to lower mean(negative)
t.test(top10$negative ~ (top10$homethatweek_num > median(top10$homethatweek_num, na.rm = TRUE)),
       data = top10, alternative = "greater")

##
## Welch Two Sample t-test
##
## data: top10$negative by top10$homethatweek_num > median(top10$homethatweek_num, na.rm = TRUE)
## t = 0.95468, df = 36.294, p-value = 0.173
## alternative hypothesis: true difference in means between group FALSE and group TRUE is greater than 0
## 95 percent confidence interval:

```

```

## -0.9666193      Inf
## sample estimates:
## mean in group FALSE  mean in group TRUE
##          15.86944          14.61093

# we fail to reject the null hypothesis, as p-value is above 0.05

#H0: mean(exhaustion) does not differ from homethatweek_num
#H1: higher homethatweek_num leads to lower mean(exhaustion)
t.test(top10$exhaustion ~ (top10$homethatweek_num > median(top10$homethatweek_num, na.rm = TRUE)),
       data = top10, alternative = "greater")

##
## Welch Two Sample t-test
##
## data:  top10$exhaustion by top10$homethatweek_num > median(top10$homethatweek_num, na.rm = TRUE)
## t = -0.95967, df = 35.539, p-value = 0.8281
## alternative hypothesis: true difference in means between group FALSE and group TRUE is greater than 0
## 95 percent confidence interval:
## -4.45019      Inf
## sample estimates:
## mean in group FALSE  mean in group TRUE
##          6.566667          8.179143

```